

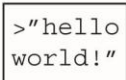


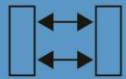
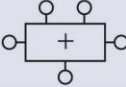
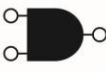
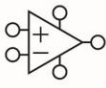

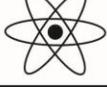
Chapter 7

Digital Design and Computer Architecture, 2nd Edition

David Money Harris and Sarah L. Harris

Chapter 7 :: Topics

- Introduction
- Performance Analysis
- Single-Cycle Processor
- Pipelined Processor
- Exceptions
- Advanced Microarchitecture

Application Software	
Operating Systems	
Architecture	
Micro-architecture	
Logic	
Digital Circuits	
Analog Circuits	
Devices	
Physics	

Introduction

- **Microarchitecture:** the implementation of an architecture in hardware
- **Processor:**
 - **Datapath:** functional blocks
 - **Control:** control signals

Application Software	programs
Operating Systems	device drivers
Architecture	instructions registers
Micro-architecture	datapaths controllers
Logic	adders memories
Digital Circuits	AND gates NOT gates
Analog Circuits	amplifiers filters
Devices	transistors diodes
Physics	electrons



Microarchitecture

- Multiple implementations for a single architecture:
 - **Single-cycle:** Each instruction executes in a single cycle
 - **Multicycle:** Instructions are broken into series of shorter steps. Each instruction executes in n cycles, where n varies according to the instr.
 - **Pipelined:** Each instruction broken up into series of steps & multiple instructions execute at once (Note: AMD and Intel pipelines are different, for the same IA-32 architecture (a.k.a. x86 ISA))

Processor Performance

- Program execution time

Execution Time = (#instructions)(cycles/instruction)(seconds/cycle)

- Definitions:
 - IC: Instruction Count (= #instructions)
 - CPI: Cycles/Instruction
 - clock period: seconds/cycle
 - IPC: Instructions/Cycle (= 1/CPI)
- Challenge is to satisfy constraints of:
 - Cost
 - Power
 - Performance

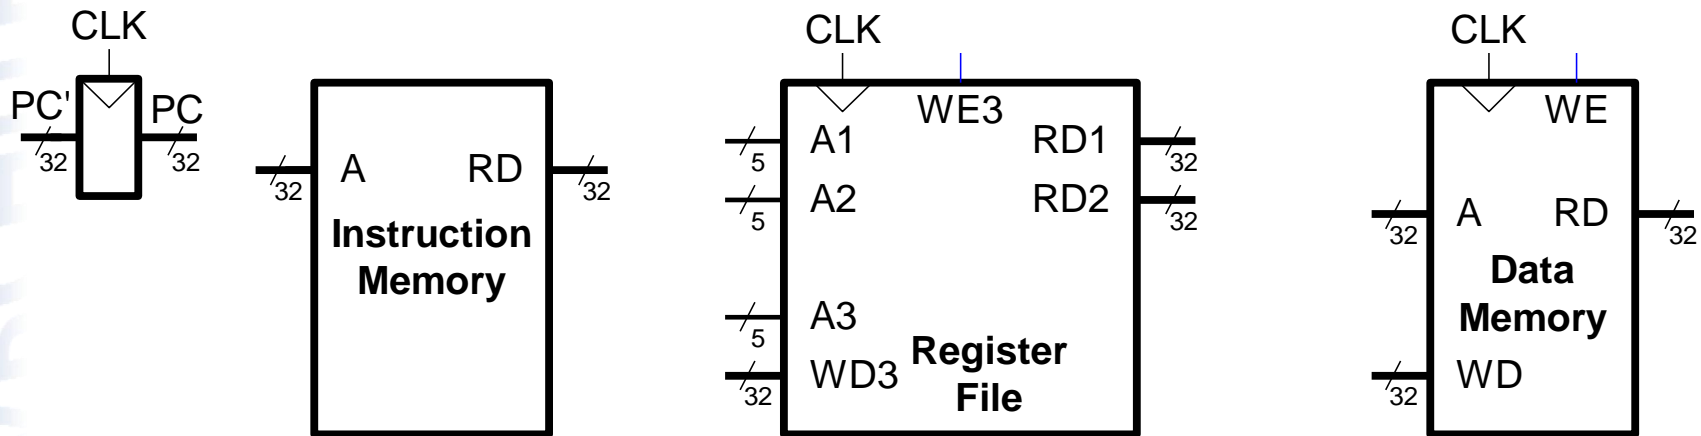
MIPS Processor

- Consider subset of MIPS instructions:
 - R-type instructions: `and`, `or`, `add`, `sub`, `sllt`
 - Memory instructions: `lw`, `sw`
 - Branch instructions: `beq`

Architectural State

- Determines everything about a processor:
 - PC and special registers
 - Register File
 - Memory

MIPS State Elements



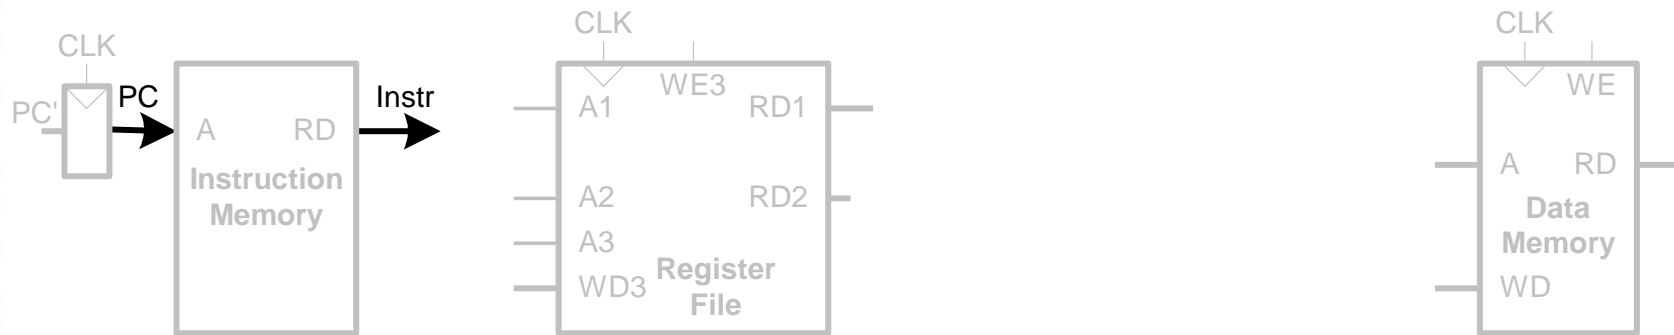
Plus the HI and LO registers

Single-Cycle MIPS Processor

- Datapath—design it 1st, to make the instruction actions possible
- Control—design it 2nd, to make them happen


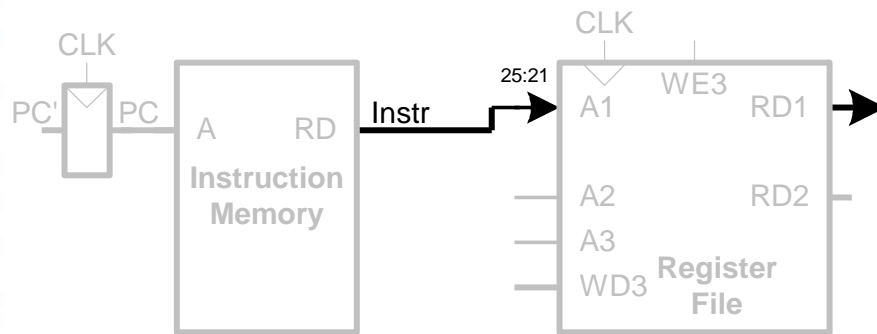
Single-Cycle Datapath: 1_w fetch

STEP 1: Fetch instruction $IM[PC]$



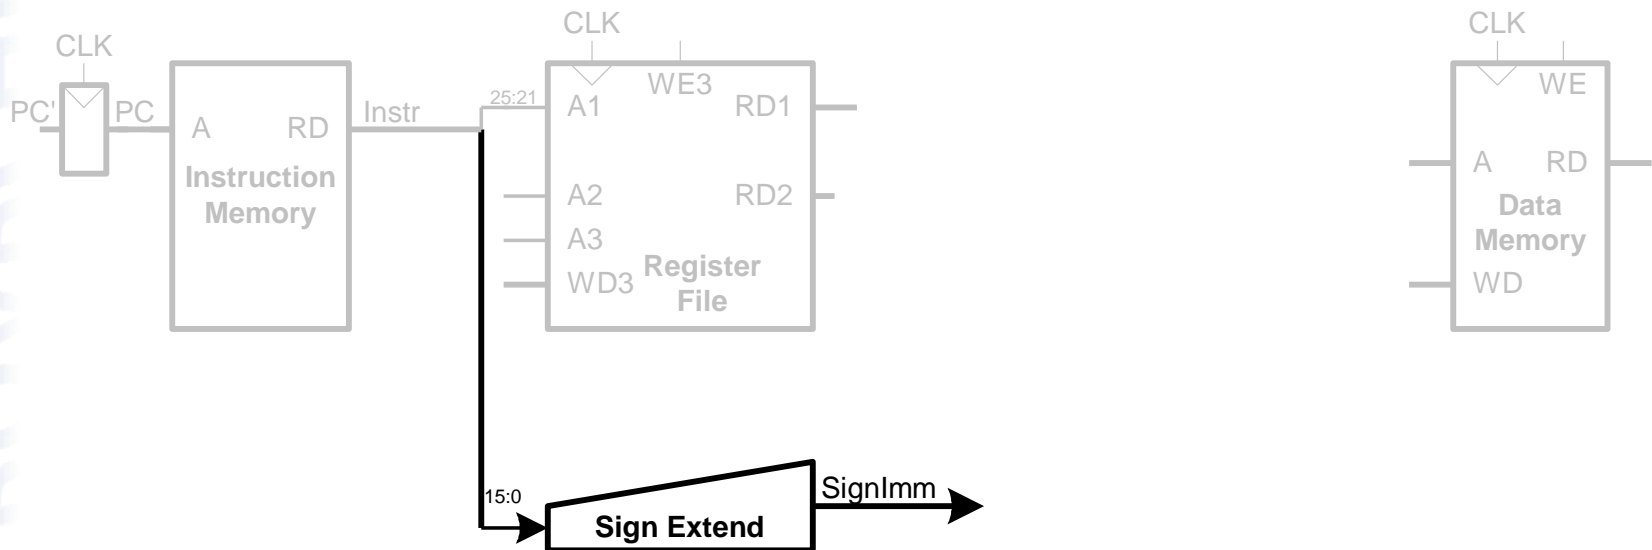
Single-Cycle Datapath: 1_w Register Read

STEP 2: Read source operands from RF
RF[rs] or RF[Instr(25:21)]



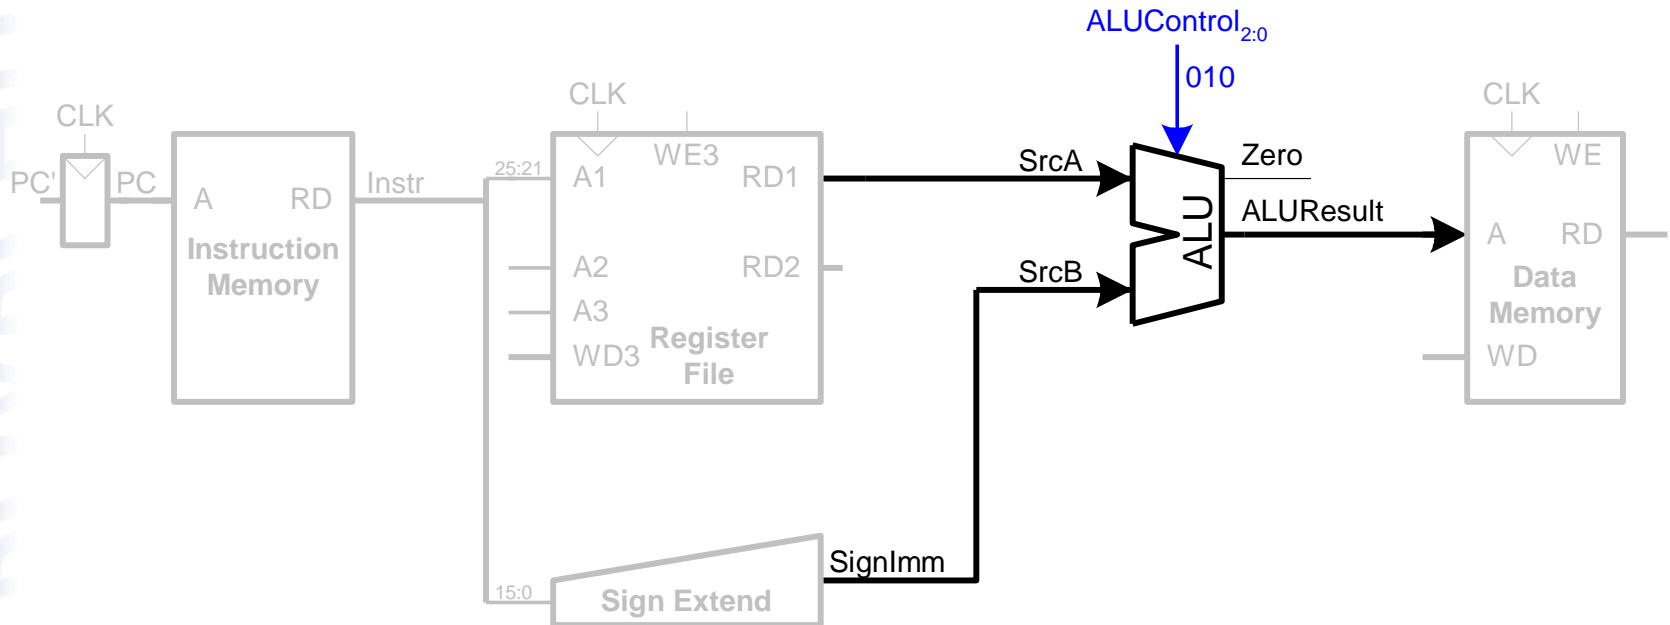
Single-Cycle Datapath: 1_w Immediate

STEP 3: Sign-extend the immediate $\text{SignExt}(\text{immed})$



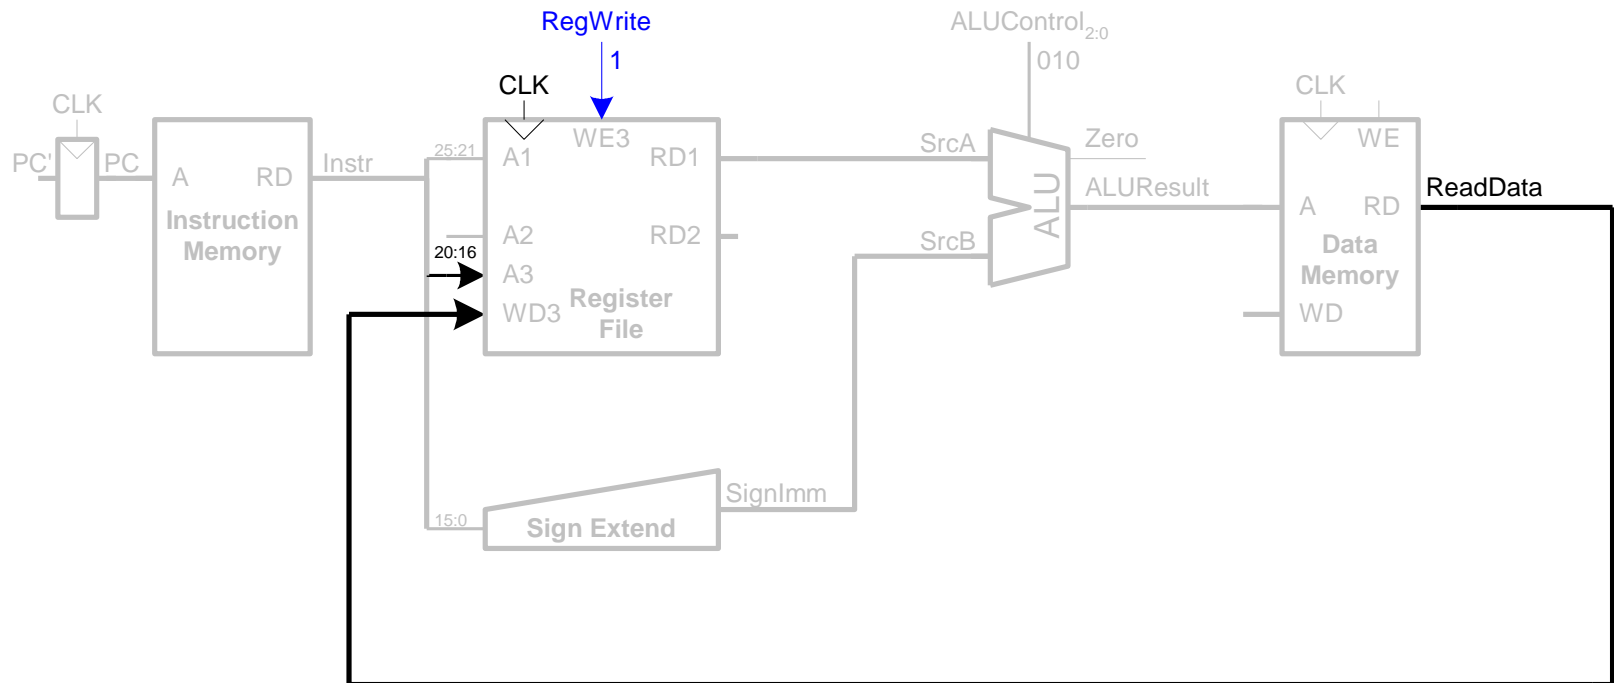
Single-Cycle Datapath: lw address

STEP 4: Compute the memory address
$$\text{addr} = \text{RF}[\text{rs}] + \text{SignExt}(\text{immed})$$



Single-Cycle Datapath: l_w Memory Read

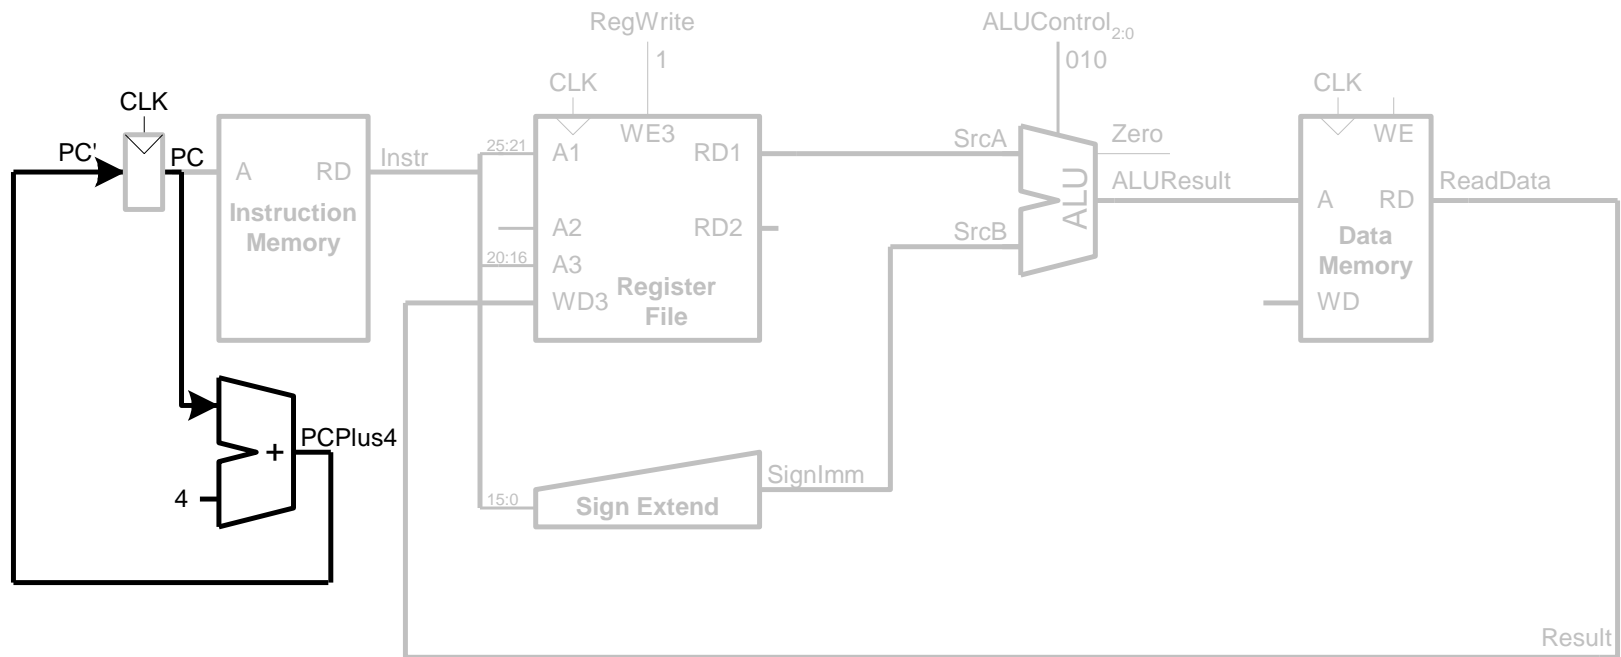
- STEP 5:** Read data from memory and write it back to register file: $RF[rt] \leftarrow DM[addr]$



Single-Cycle Datapath: 1_w PC Increment

STEP 6: Determine address of next instruction

$$PC \leftarrow PC + 4$$



Full RTL Expression for \perp_w

IM[PC]

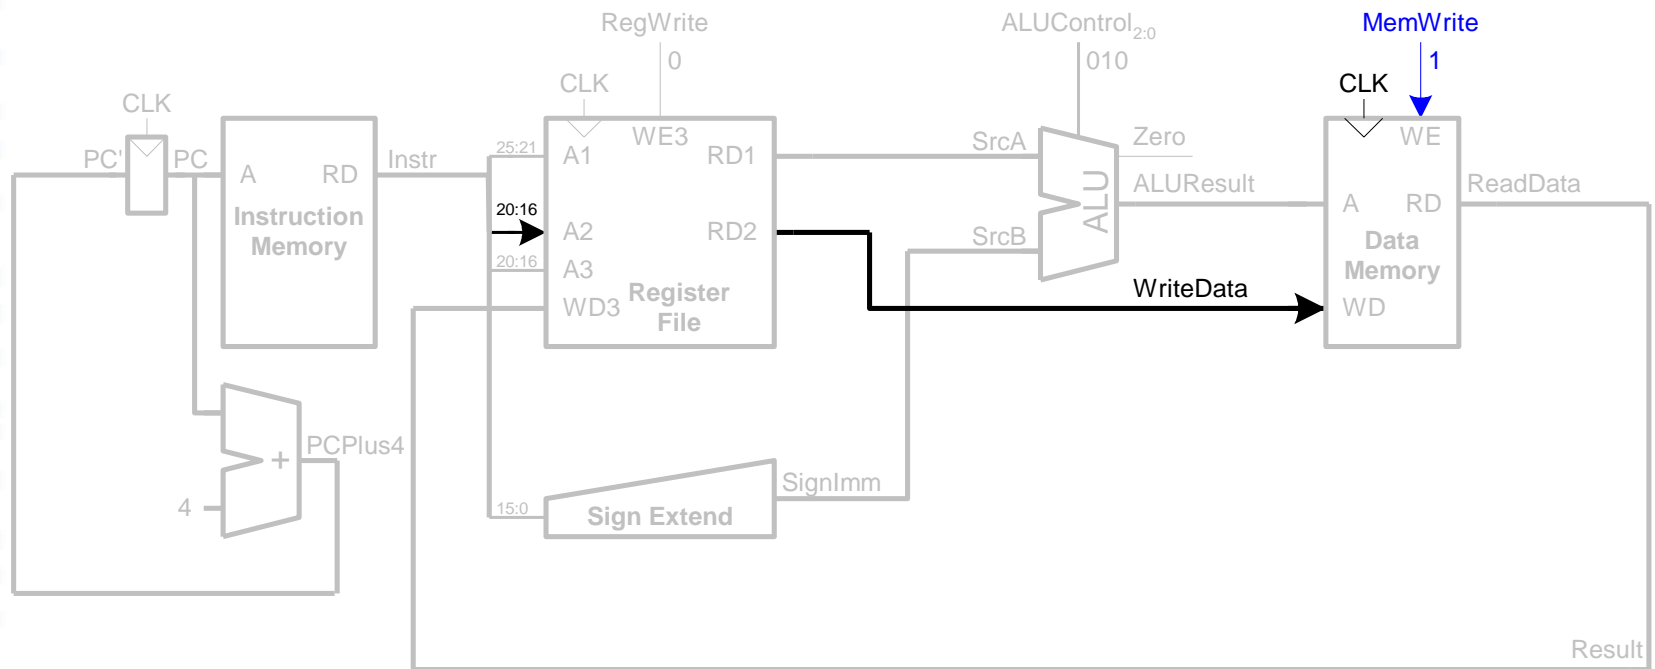
32 bit + 32bit(16 extended)

$RF[rt] \leftarrow DM[RF[rs] + \text{SignExt(immed)}]$

$PC \leftarrow PC + 4$

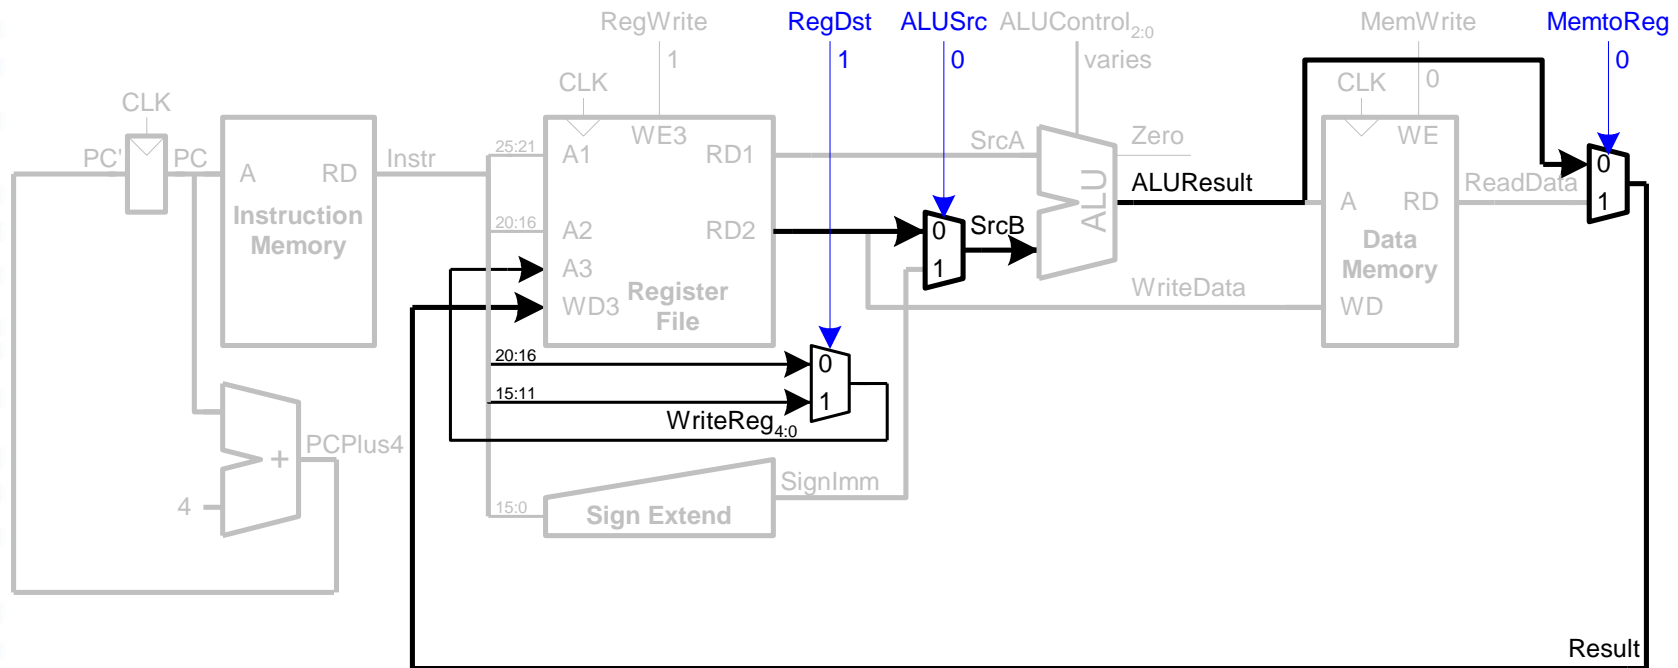
Single-Cycle Datapath: sw

Write data in rt to memory: $DM[addr] \leftarrow RF[rt]$



Single-Cycle Datapath: R-Type

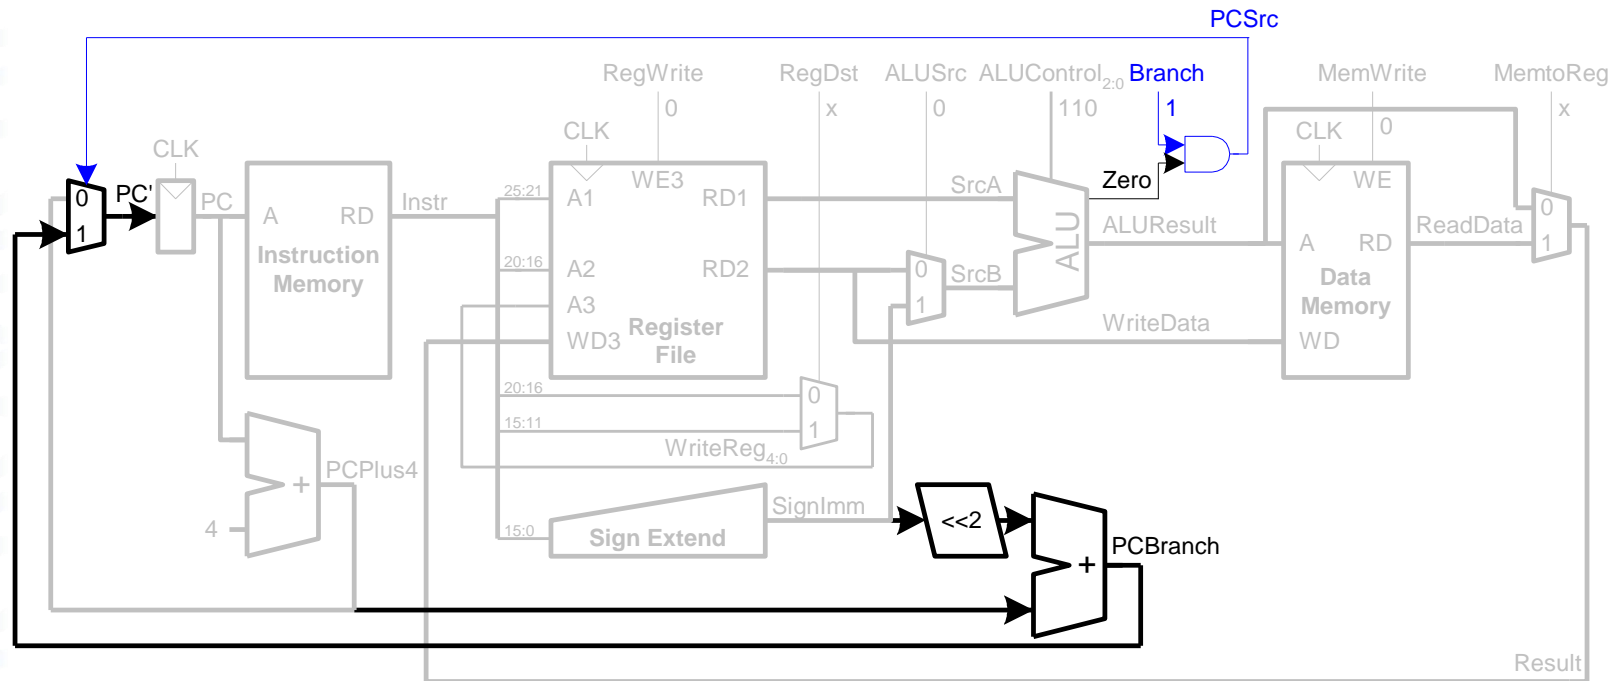
- Read from rs and rt
- Write $ALUResult$ to register file
- Write to rd (instead of rt) $RF[rd] \leftarrow RF[rs] \text{ op } RF[rt]$



Single-Cycle Datapath: beq

- Determine whether values in rs and rt are equal
- Calculate branch target address:

$$BTA = PC + 4 + \text{SignExt}(\text{immed}) \ll 2 \quad \# \ll 2 = 4x$$



RTL Expression for beq

IM[PC]

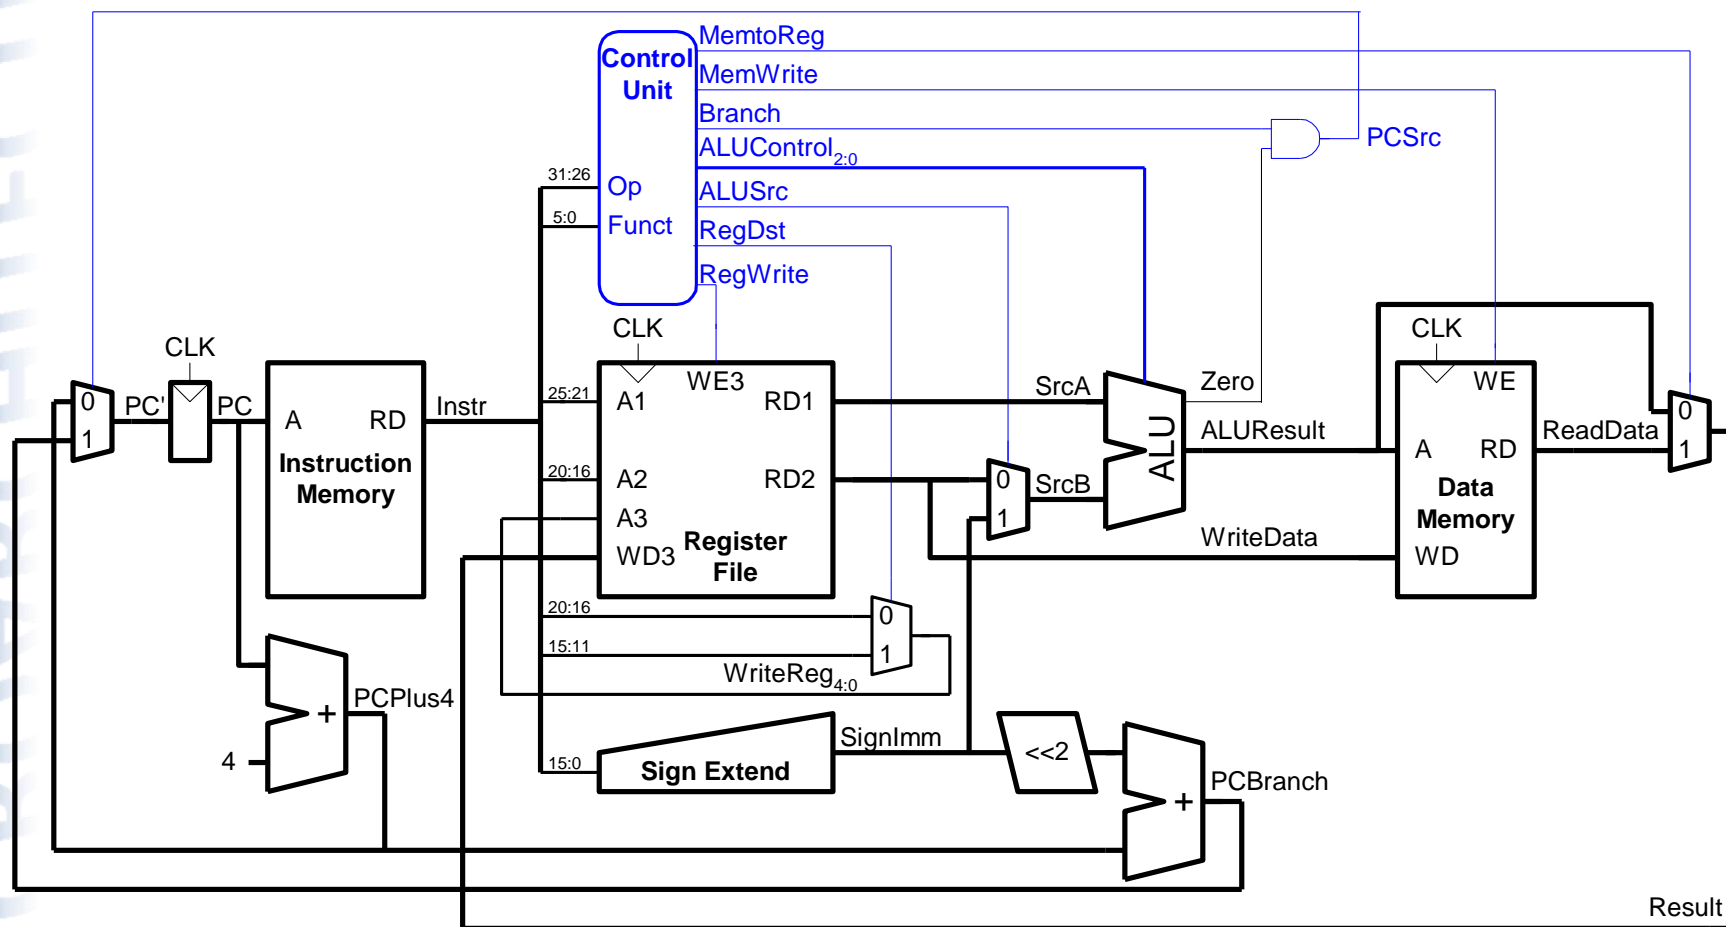
if (RF[rs] - RF[rt] == 0)

PC \leftarrow BTA

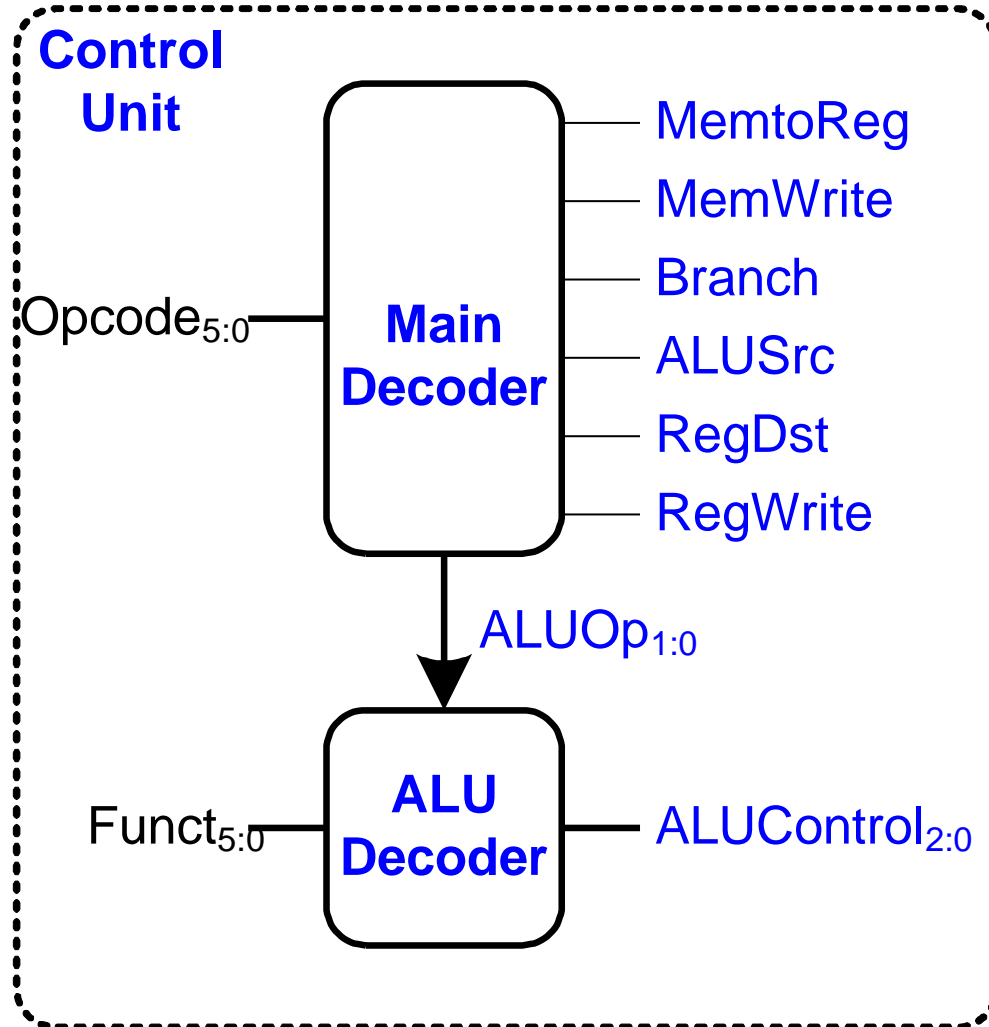
else

PC \leftarrow PC + 4


Single-Cycle Processor



Single-Cycle Control




Review: ALU



$F_{2:0}$	Function
000	$A \& B$
001	$A B$
010	$A + B$
011	not used
100	$A \& \sim B$
101	$A \sim B$
110	$A - B$
111	SLT

Review: ALU



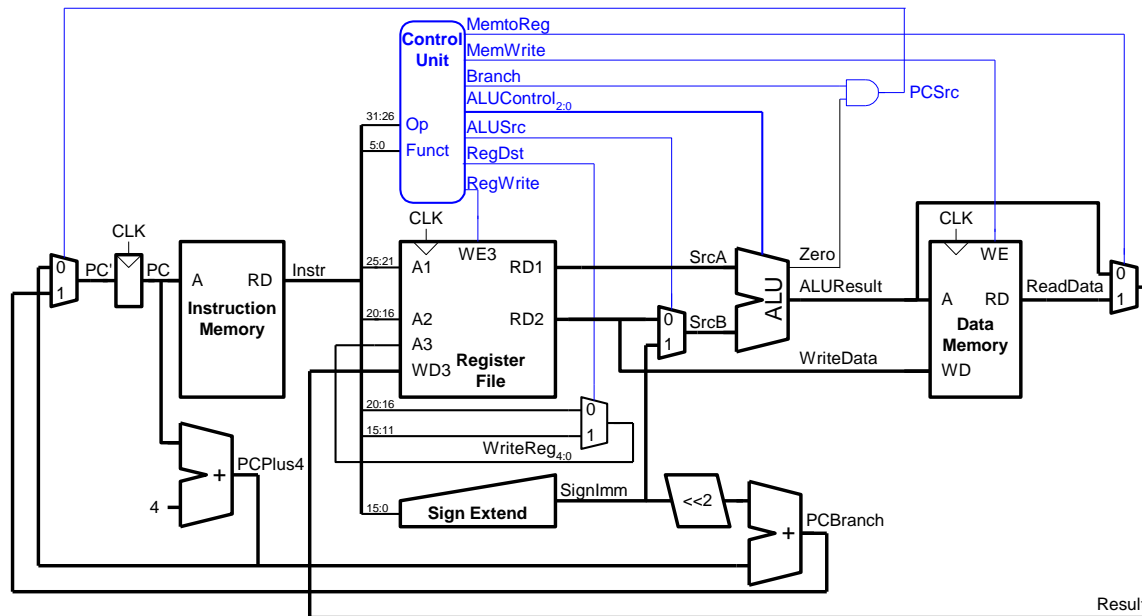
Control Unit: ALU Decoder

ALUOp _{1:0}	Meaning
00	Add (for lw, sw)
01	Subtract (for beq)
10	Look at funct (R-type)
11	Not Used

ALUOp _{1:0}	funct	ALUControl _{2:0}
00	X	010 (Add)
X1	X	110 (Subtract)
1X	100000 (add)	010 (Add)
1X	100010 (sub)	110 (Subtract)
1X	100100 (and)	000 (And)
1X	100101 (or)	001 (Or)
1X	101010 (slt)	111 (SLT)

Control Unit: Main Decoder

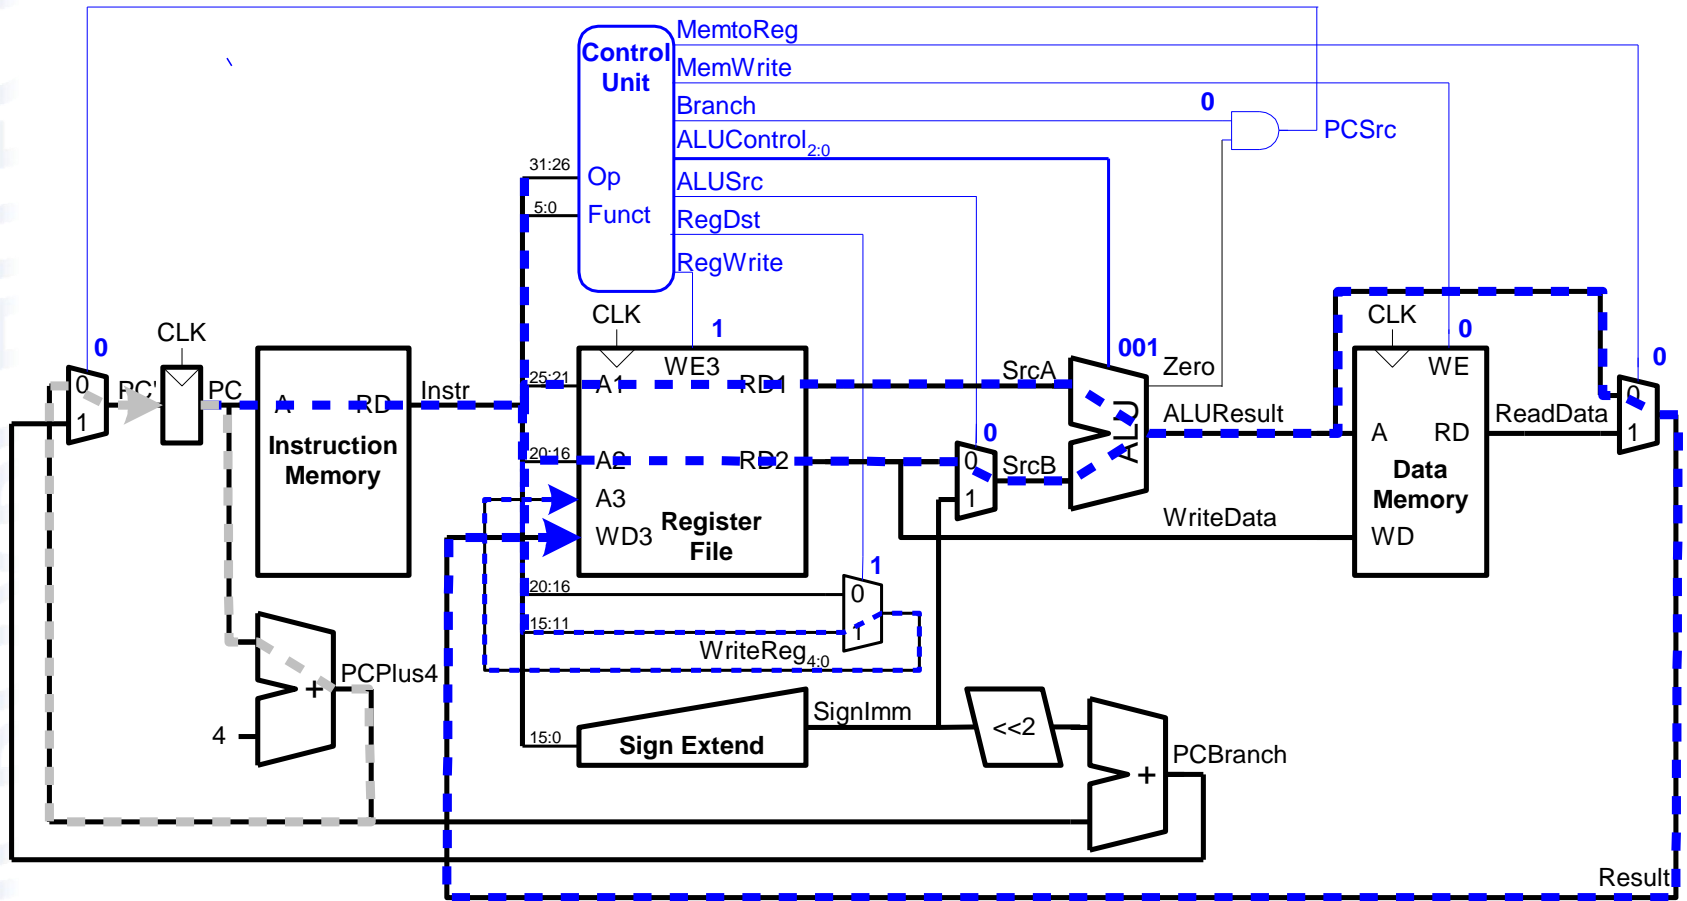
Instruction	Op _{5:0}	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp _{1:0}
R-type	000000							
lw	100011							
sw	101011							
beq	000100							



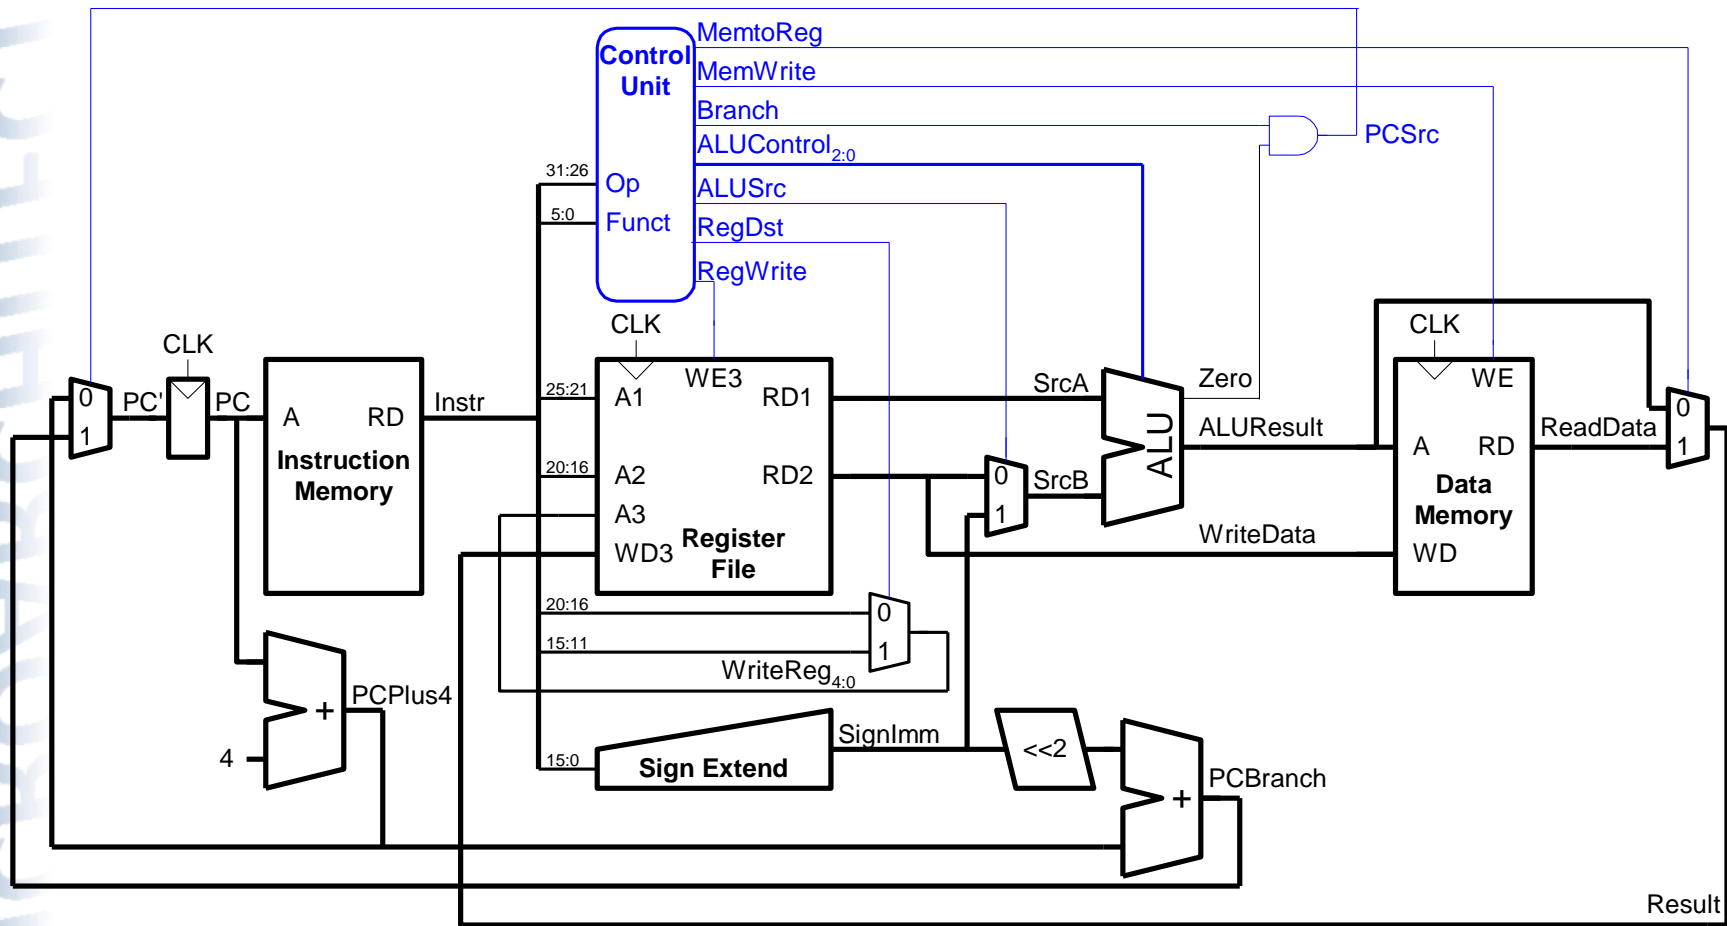
Control Unit: Main Decoder

Instruction	Op _{5:0}	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp _{1:0}
R-type	000000	1	1	0	0	0	0	10
lw	100011	1	0	1	0	0	1	00
sw	101011	0	X	1	0	1	X	00
beq	000100	0	X	0	1	0	X	01

Single-Cycle Datapath: or



Extended Functionality: addi



No change to datapath

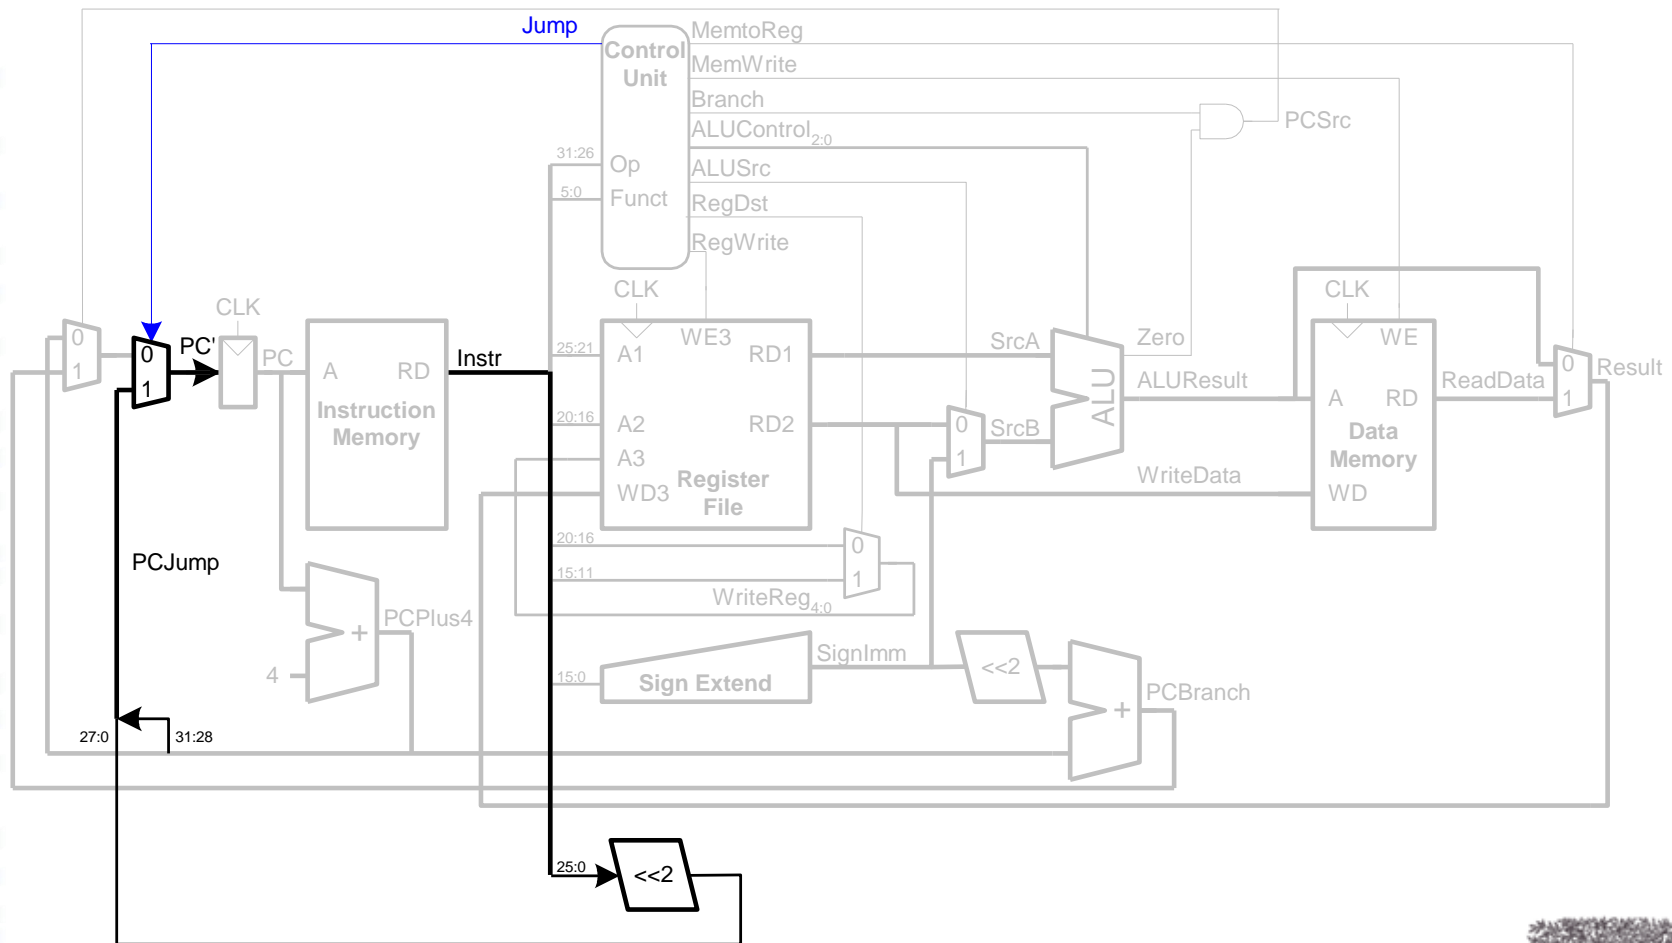
Main Decoder table: addi

Instruction	Op _{5:0}	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp _{1:0}
R-type	000000	1	1	0	0	0	0	10
lw	100011	1	0	1	0	0	1	00
sw	101011	0	X	1	0	1	X	00
beq	000100	0	X	0	1	0	X	01
addi	001000							

Main Decoder table: addi

Instruction	Op _{5:0}	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp _{1:0}
R-type	000000	1	1	0	0	0	0	10
lw	100011	1	0	1	0	0	1	00
sw	101011	0	X	1	0	1	X	00
beq	000100	0	X	0	1	0	X	01
addi	001000	1	0	1	0	0	0	00

Extended Functionality: j



Main Decoder table: j

Instruction	Op _{5:0}	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp _{1:0}	Jump
R-type	000000	1	1	0	0	0	0	10	0
lw	100011	1	0	1	0	0	1	00	0
sw	101011	0	X	1	0	1	X	00	0
beq	000100	0	X	0	1	0	X	01	0
j	000010								

Main Decoder table: j

Instruction	Op _{5:0}	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp _{1:0}	Jump
R-type	000000	1	1	0	0	0	0	10	0
lw	100011	1	0	1	0	0	1	00	0
sw	101011	0	X	1	0	1	X	00	0
beq	000100	0	X	0	1	0	X	01	0
j	000010	0	X	X	X	0	X	XX	1

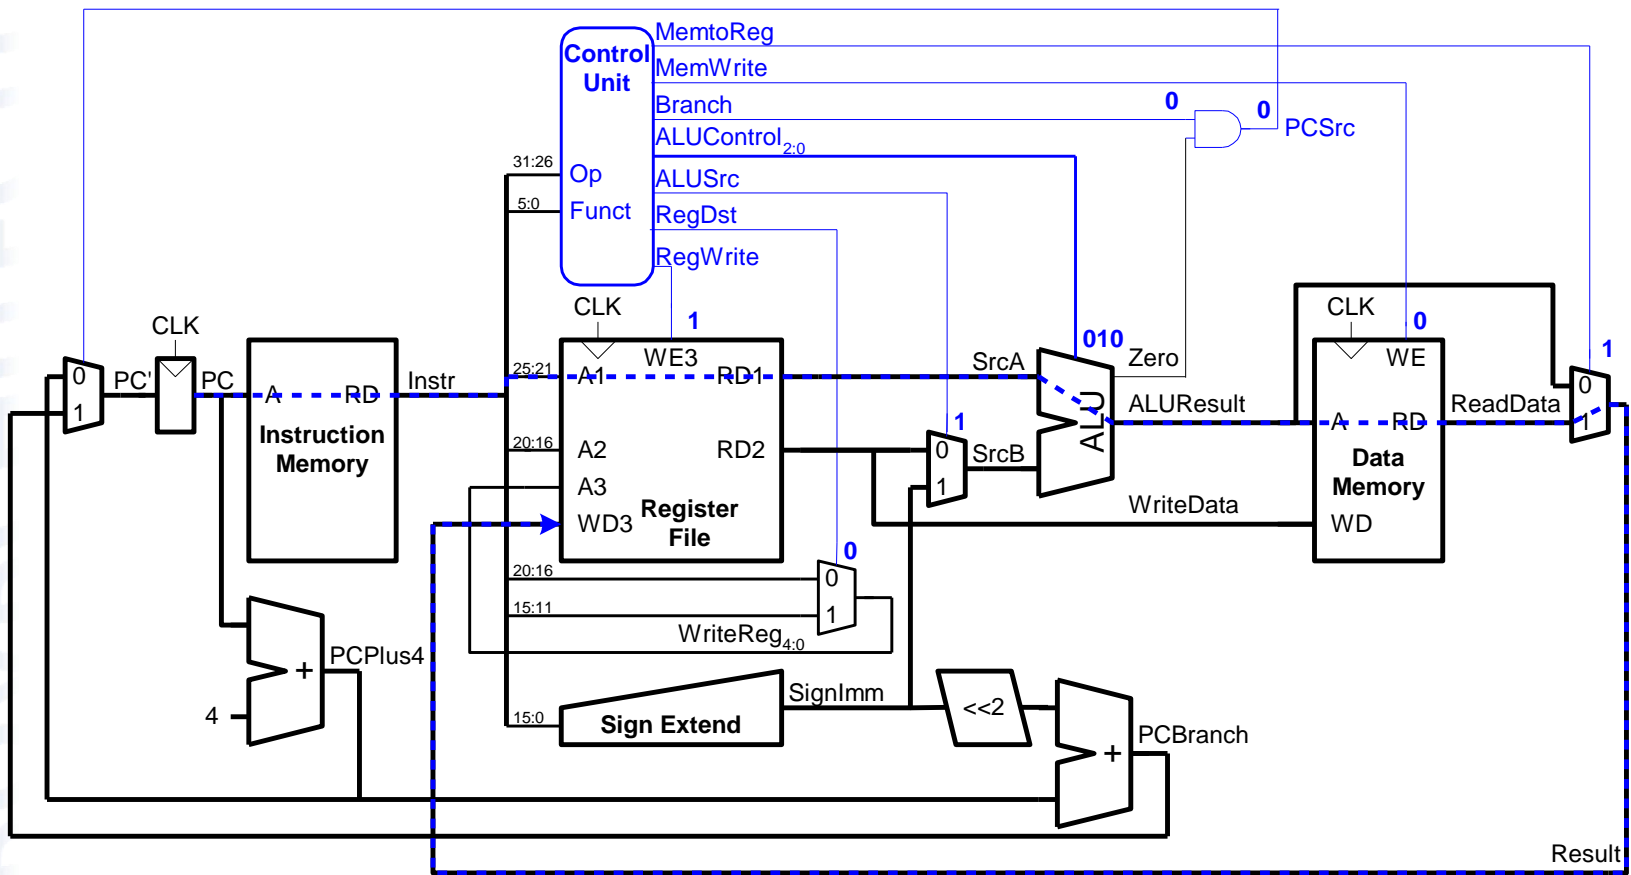
Review: Processor Performance

Program Execution Time

$$= (\text{\#instructions})(\text{cycles/instruction})(\text{seconds/cycle})$$

$$= IC \times CPI \times T_c$$

Single-Cycle Performance



T_C limited by critical path (1w)

Single-Cycle Performance

- Single-cycle critical path:

$$T_c = t_{pcq_PC} + t_{mem} + \max(t_{RFread}, t_{sext} + t_{mux}) + t_{ALU} + t_{mem} + t_{mux} + t_{RFsetup}$$

- Typically, limiting paths are:

- memory, ALU, register file

- $T_c = t_{pcq_PC} + 2t_{mem} + t_{RFread} + t_{mux} + t_{ALU} + t_{RFsetup}$



Single-Cycle Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	t_{pcq_PC}	30
Register setup	t_{setup}	20
Multiplexer	t_{mux}	25
ALU	t_{ALU}	200
Memory read	t_{mem}	250
Register file read	t_{RFread}	150
Register file setup	$t_{RFsetup}$	20

$$T_c = ?$$

Single-Cycle Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	t_{pcq_PC}	30
Register setup	t_{setup}	20
Multiplexer	t_{mux}	25
ALU	t_{ALU}	200
Memory read	t_{mem}	250
Register file read	t_{RFread}	150
Register file setup	$t_{RFsetup}$	20

$$\begin{aligned}T_c &= t_{pcq_PC} + 2t_{mem} + t_{RFread} + t_{mux} + t_{ALU} + t_{RFsetup} \\&= [30 + 2(250) + 150 + 25 + 200 + 20] \text{ ps} \\&= 925 \text{ ps} \quad [f_{clk} = 1/0.925 \text{ GHz} = 1.08 \text{ GHz}]\end{aligned}$$

Single-Cycle Performance Example

Program with IC = 100 billion instructions:

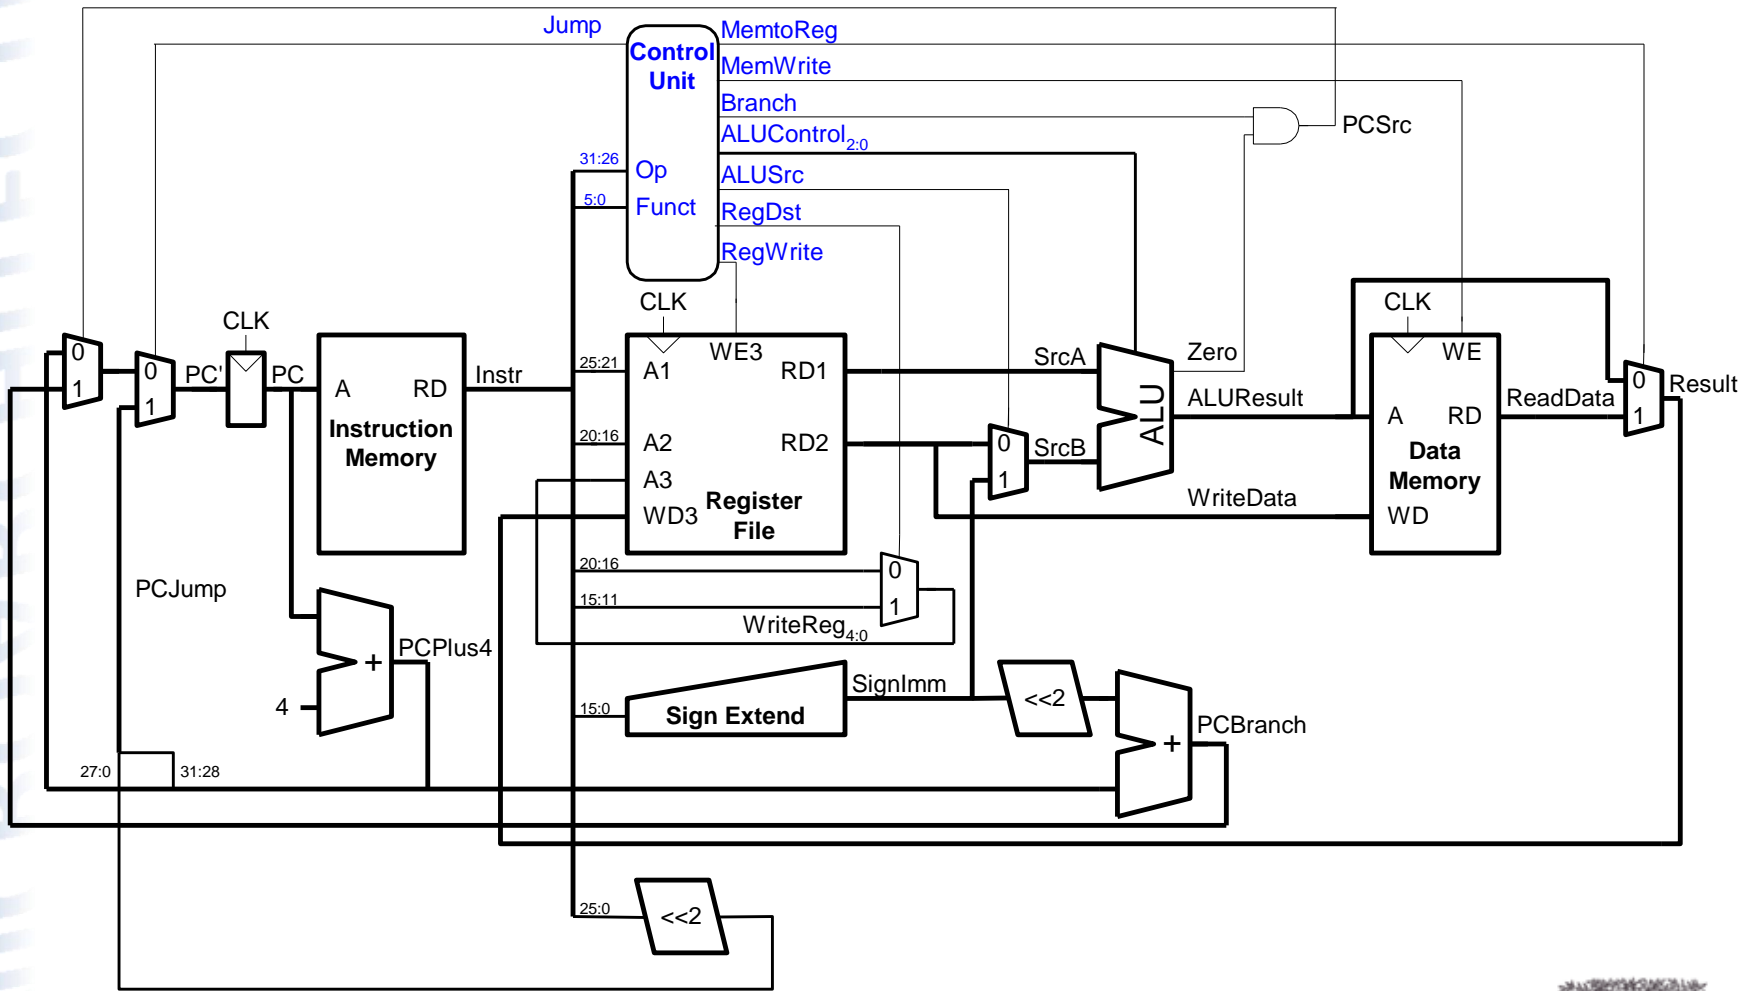
$$\begin{aligned}\text{Execution Time} &= \text{IC} \times \text{CPI} \times T_C \\ &= (100 \times 10^9)(1)(925 \times 10^{-12} \text{ s}) \\ &= \mathbf{92.5 \text{ seconds}}\end{aligned}$$

Evaluation of Single-Cycle Processor

Pros and cons of single-cycle implementation:

- + simple design
- + 1 cycle per every instruction
- slow cycle time
 - limited by longest instruction (1_w)
- HW: 2 adders + ALU; 2 memories

Review: Single-Cycle Processor



System Verilog Model

```
//-----  
// by David Harris, in Chapter 7 of DDCA 2nd ed. textbook  
// Top level system of Fig 7.59 including MIPS and memories  
//-----  
  
module top (input  logic      clk, reset,  
            output logic[31:0] writedata, dataadr,  
            output logic      memwrite);  
  
    logic [31:0] pc, instr, readdata;  
  
    // instantiate processor and memories  
    mips mips (clk, reset, pc, instr, memwrite, dataadr, writedata, readdata);  
    imem imem (pc[7:2], instr);  
    dmem dmem (clk, memwrite, dataadr, writedata, readdata);  
  
endmodule
```

Sys. Verilog Model of Data Memory

```
//-----  
// External data memory used by MIPS single-cycle processor  
//-----  
module dmem (input logic      clk, we,  
              input logic[31:0] a, wd,  
              output logic[31:0] rd);  
  
logic [31:0] RAM[63:0];  
  
assign rd = RAM[a[31:2]]; // word-aligned read  
  
always_ff @(posedge clk)  
    if (we)  
        RAM[a[31:2]] <= wd; // word-aligned write  
endmodule
```

Sys. Verilog Model of Instr. Memory

```
module imem (input logic[5:0] addr, output reg [31:0] instr);

// imem is modeled as a lookup table, a stored-program byte-addressable ROM
always_comb
    case ({addr,2'b00})
//          address          instruction
//          -----          -
        8'h00: instr = 32'h20020005;
        8'h04: instr = 32'h2003000c;
        8'h08: instr = 32'h2067fff7;
        8'h0c: instr = 32'h00e22025;
        8'h10: instr = 32'h00642824;
        8'h14: instr = 32'h00a42820;
        8'h18: instr = 32'h10a7000a;
        8'h1c: instr = 32'h0064202a;
        8'h20: instr = 32'h10800001;
        default: instr = {32{1'bx}}; // unknown instruction
    endcase
endmodule
```



Alternate Model of Instr. Memory

```
module imem (input  logic[5:0] addr,
             output logic[31:0] instr);

logic [31:0] RAM[63:0];

// imem is RAM, loaded from memfile.dat file with hex values at startup
initial
begin
    $readmemh("memfile.dat", RAM);
end

assign instr = RAM[addr]; // instr at RAM[addr] is read out

endmodule

// imem can be created with CoreGen for Xilinx synthesis
```

Sys. Verilog Model of MIPS processor

```
// single-cycle MIPS processor
module mips (input logic      clk, reset,
             output logic[31:0] pc,
             input logic[31:0] instr,
             output logic      memwrite,
             output logic[31:0] aluout, writedata,
             input logic[31:0] readdata);

    logic      memtoreg, pcsrc, zero, alusrc, regdst, regwrite, jump;
    logic [2:0] alucontrol;

    controller c (instr[31:26], instr[5:0], zero, memtoreg, memwrite, pcsrc,
                  alusrc, regdst, regwrite, jump, alucontrol);

    datapath dp (clk, reset, memtoreg, pcsrc, alusrc, regdst, regwrite, jump,
                 alucontrol, zero, pc, instr, aluout, writedata, readdata);

endmodule
```

Sys. Verilog Model of Controller

```
module controller ( input logic[5:0] op, funct,
                    input logic      zero,
                    output logic      memtoreg, memwrite,
                    output logic      pcsrc, alusrc,
                    output logic      regdst, regwrite,
                    output logic      jump,
                    output logic[2:0] alucontrol);

    logic [1:0] aluop;
    logic      branch;

    maindec md (op, memtoreg, memwrite, branch, alusrc, regdst, regwrite,
                jump, aluop);

    aludec ad (funct, aluop, alucontrol);

    assign pcsrc = branch & zero;

endmodule
```


Sys. Verilog Model of Main Decoder

```
module maindec (input  logic[5:0] op,
                output logic memtoreg, memwrite, branch,
                output logic alusrc, regdst, regwrite, jump,
                output logic[1:0] aluop );
    logic [8:0] controls;

    assign {regwrite, regdst, alusrc, branch, memwrite,
            memtoreg, aluop, jump} = controls;

    always_comb
    case(op)
        6'b000000: controls = 9'b1100000010; //Rtype
        6'b100011: controls = 9'b101001000; //LW
        6'b101011: controls = 9'b001010000; //SW
        6'b000100: controls = 9'b000100001; //BEQ
        6'b001000: controls = 9'b101000000; //ADDI
        6'b000010: controls = 9'b0000000100; //J
        default:   controls = 9'bxxxxxxxx; //illegal op
    endcase
endmodule
```

Sys. Verilog Model of ALU Decoder

```
module aludec (input  logic[5:0] funct,
               input  logic[1:0] aluop,
               output logic [2:0] alucontrol);

always_comb
  case(aluop)
    2'b00: alucontrol = 3'b010; // add (for lw/sw/addi)
    2'b01: alucontrol = 3'b110; // sub (for beq)
    default: case(funct)          // R-TYPE instructions
      6'b100000: alucontrol = 3'b010; // ADD
      6'b100010: alucontrol = 3'b110; // SUB
      6'b100100: alucontrol = 3'b000; // AND
      6'b100101: alucontrol = 3'b001; // OR
      6'b101010: alucontrol = 3'b111; // SLT
      default:   alucontrol = 3'bxxx; // ???
    endcase
  endcase
endmodule
```



Sys. Verilog Model of Datapath

```
module datapath (input  logic clk, reset, memtoreg, pcsrc, alusrc, regdst,
                input  logic regwrite, jump, input  logic[2:0] alucontrol,
                output logic zero, output logic[31:0] pc,
                input  logic[31:0] instr,
                output logic[31:0] aluout, writedata,
                input  logic[31:0] readdata);

    logic [4:0] writereg;
    logic [31:0] pcnext, pcnextbr, pcplus4, pcbranch;
    logic [31:0] signimm, signimmsh, srca, srcb, result;

    // next PC logic
    flopr #(32) pcreg(clk, reset, pcnext, pc);
    adder      pcadd1(pc, 32'b100, pcplus4);
    sl2        immsh(signimm, signimmsh);
    adder      pcadd2(pcplus4, signimmsh, pcbranch);
    mux2 #(32) pcbrmux(pcplus4, pcbranch, pcsrc,
                     pcnextbr);
    mux2 #(32) pcmux(pcnextbr, {pcplus4[31:28],
                              instr[25:0], 2'b00}, jump, pcnext);
```

Sys. Verilog Model of Datapath

// register file logic

```
regfile  rf (clk, regwrite, instr[25:21], instr[20:16], writereg,  
            result, srca, writedata);
```

```
mux2 #(5)  wrmux (instr[20:16], instr[15:11], regdst, writereg);
```

```
mux2 #(32) resmux (aluout, readdata, memtoreg, result);
```

```
signext    se (instr[15:0], signimm);
```

// ALU logic

```
mux2 #(32) srcbmux (writedata, signimm, alusrc, srcb);
```

```
alu        alu (srca, srcb, alucontrol, aluout, zero);
```

```
endmodule
```

Sys. Verilog Model of Register File

```
module regfile (input  logic clk, we3,
                input  logic[4:0] ra1, ra2, wa3,
                input  logic[31:0] wd3,
                output logic[31:0] rd1, rd2);

    logic [31:0] rf [31:0];

    // three ported register file: read two ports combinationaly
    // write third port on rising edge of clock. Register0 hardwired to 0.

    always_ff
        if (we3)
            rf [wa3] <= wd3;

    assign rd1 = (ra1 != 0) ? rf [ra1] : 0;
    assign rd2 = (ra2 != 0) ? rf[ ra2] : 0;

endmodule
```



Sys. Verilog Models of Other Parts

```
module adder (input  logic[31:0] a, b,  
              output logic[31:0] y);  
    assign y = a + b;  
endmodule
```

```
module sl2 (input  logic[31:0] a,  
            output logic[31:0] y);  
    // shift left by 2  
    assign y = {a[29:0], 2'b00};  
endmodule
```

```
module signext (input  logic[15:0] a,  
                output logic[31:0] y);  
  
    assign y = {{16{a[15]}}, a};  
endmodule
```

Sys. Verilog for Parameterized Parts

```
module flopr #(parameter WIDTH = 8)
    (input logic clk, reset,
     input logic[WIDTH-1:0] d,
     output logic[WIDTH-1:0] q);
```

```
    always_ff@(posedge clk, posedge reset)
        if (reset) q <= 0;
        else      q <= d;
endmodule
```

```
module mux2 #(parameter WIDTH = 8)
    (input  logic[WIDTH-1:0] d0, d1,
     input  logic s,
     output logic[WIDTH-1:0] y);
    assign y = s ? d1 : d0;
endmodule
```

Review: Exceptions

- Unscheduled function call to *exception handler*
- Caused by:
 - Hardware, also called an *interrupt*, e.g. keyboard
 - Software, also called *traps*, e.g. undefined instruction
- When exception occurs, the processor:
 - Records cause of exception (Cause register)
 - Jumps to exception handler (0x80000180)
 - Returns to program (EPC register)

Example Exception

sequential circuits.]

Can we design a spiff

Figure 2.11 shows a box indicates that it is this case, the function is

KeyAccess



The network KeyServer, which is required by KeyServer controlled programs, cannot grant you permission to run this program. If you think you have received this message in error, please contact your KeyServer Administrator.

Visio.exe - Application Error



The exception unknown software exception (0xc06d007e) occurred in the application at location 0x7c81eb33.

OK

words, we say the output Y is a function of the two inputs A and B where the function performed is $A \text{ OR } B$.]

The *implementation* of the combinational circuit is independent of its functionality. Figure 2.1 and Figure 2.2 show two possible implementa-



Review: Exception Registers

- Not part of register file; in Coprocessor 0
 - Cause
 - Records cause of exception
 - Coprocessor 0 register 13
 - EPC (Exception PC)
 - Records PC where exception occurred
 - Coprocessor 0 register 14
- Move from Coprocessor 0
 - `mfc0 $t0, Cause (=mfc0 $t0, $13)`
 - Moves contents of Cause into \$t0

mfc0

010000	00000	\$t0 (8)	Cause (13)	000000000000
31:26	25:21	20:16	15:11	10:0

Review: Exception Causes

Exception	Cause
Hardware Interrupt	0x00000000
System Call	0x00000020
Breakpoint / Divide by 0	0x00000024
Undefined Instruction	0x00000028
Arithmetic Overflow	0x00000030

Extend single-cycle MIPS processor to handle last two types of exceptions

Exception RTLs

Undefined Instruction

IM[PC]

... # problem in decoding (bad op or func)

Cause $\leftarrow 40$ # = 0x28

EPC \leftarrow PC

PC \leftarrow 0x80000180 #Exception handler address

Arithmetic Overflow

IM[PC]

... # ALU operation overflows

Cause $\leftarrow 48$ # = 0x30

EPC \leftarrow PC

PC \leftarrow 0x80000180 #Exception handler address

mfc0 instruction (e.g. mfc0 \$t1, \$13)

IM[PC]

RF[rt] \leftarrow RFc0[rd]

PC \leftarrow PC + 4



MICROARCHITECTURE



Chapter 7 <61>

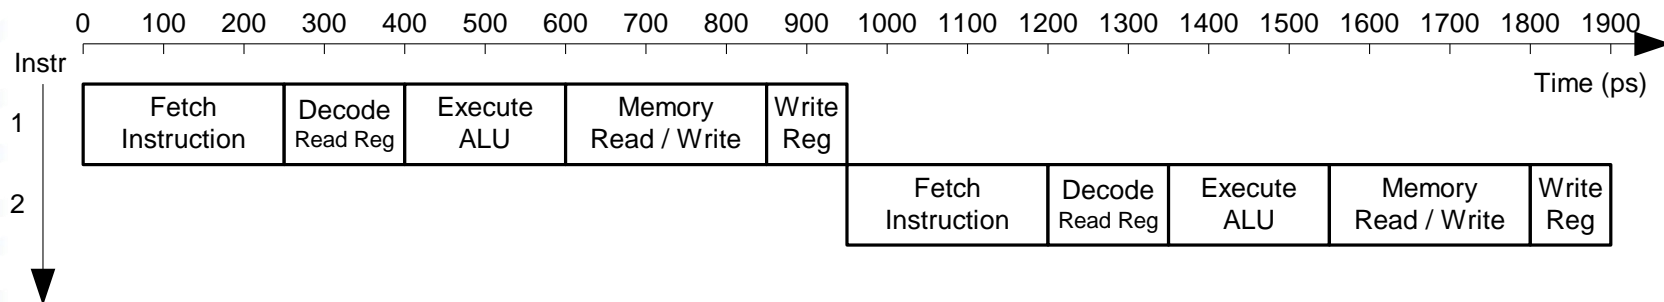


Pipelined MIPS Processor

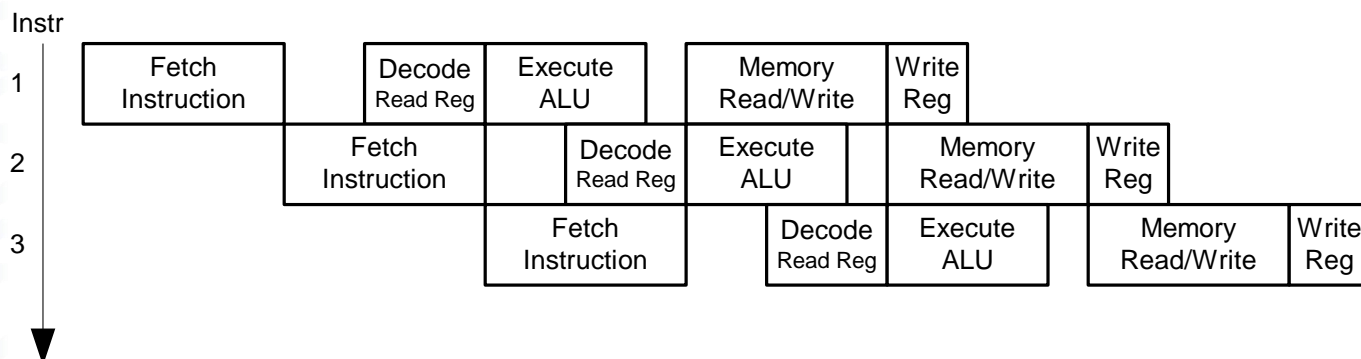
- Temporal parallelism
- Divide single-cycle processor into 5 stages:
 - Fetch
 - Decode
 - Execute
 - Memory
 - Writeback
- Add pipeline registers between stages

Single-Cycle vs. Pipelined

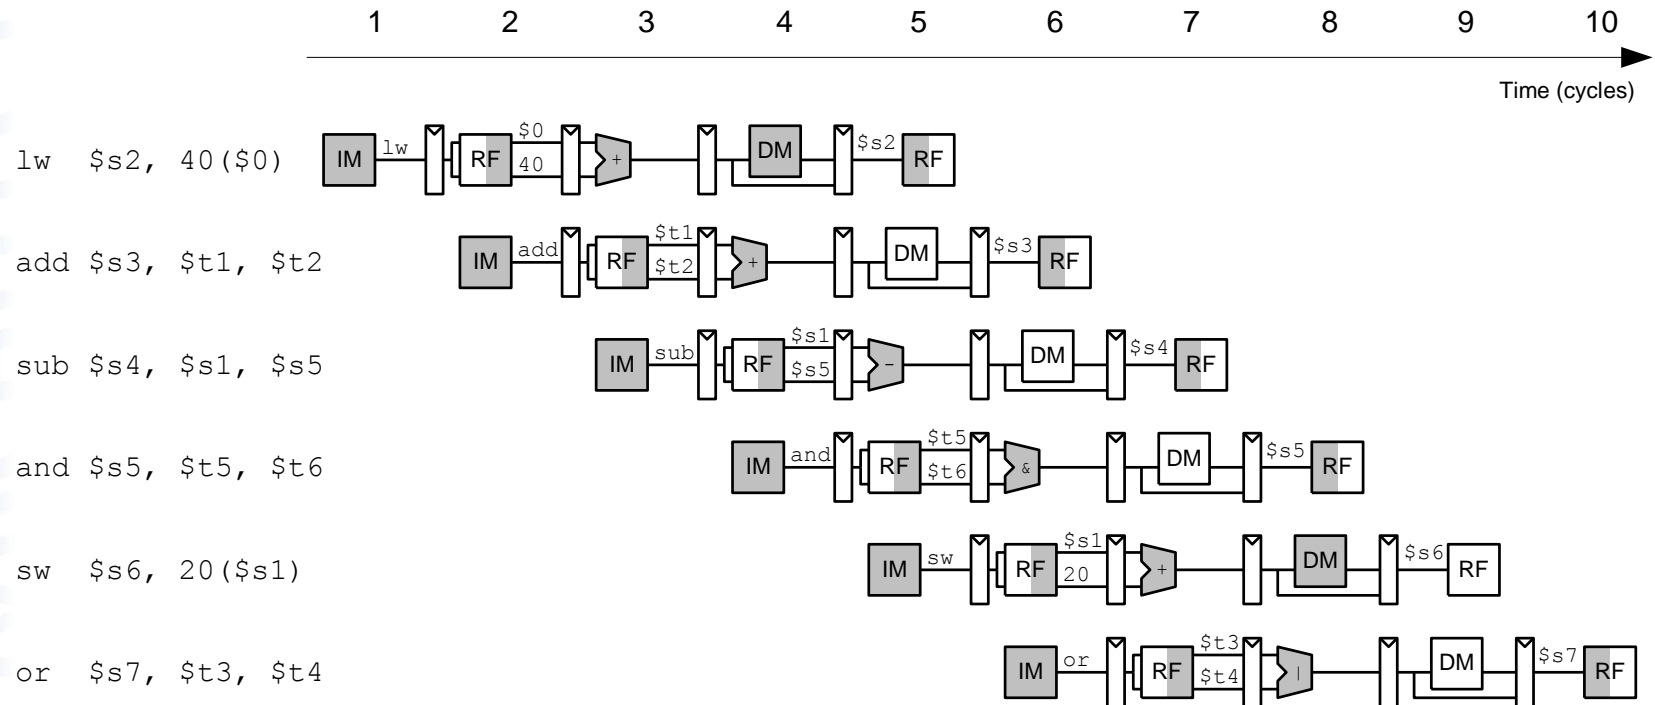
Single-Cycle



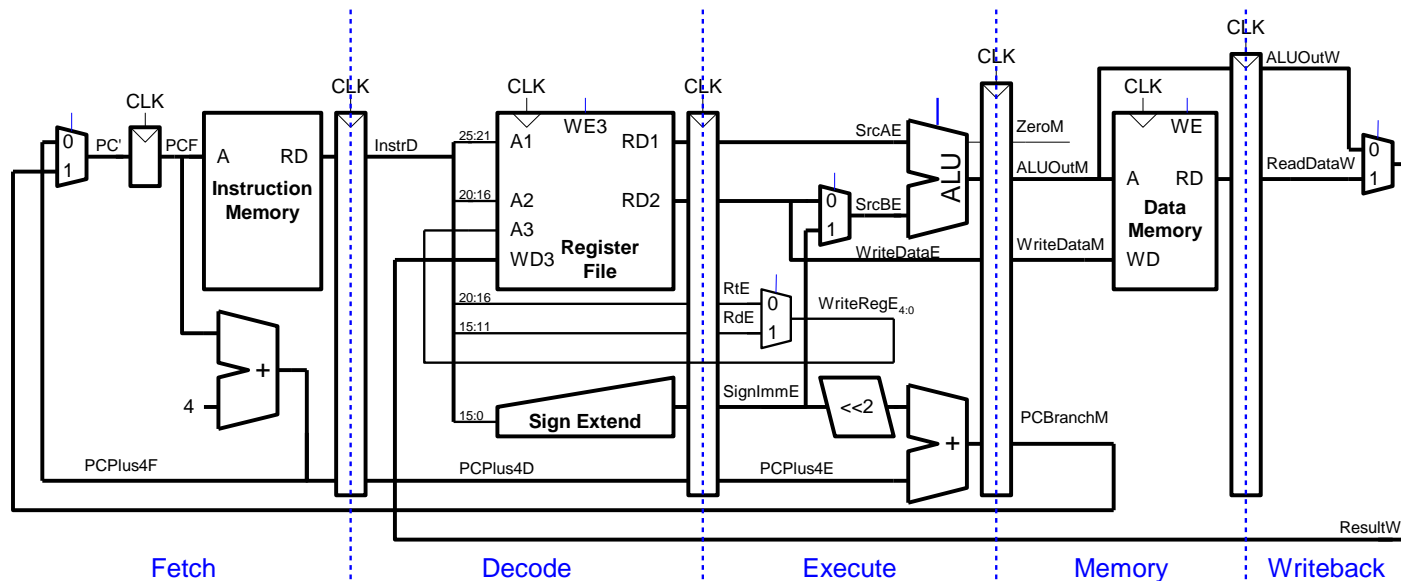
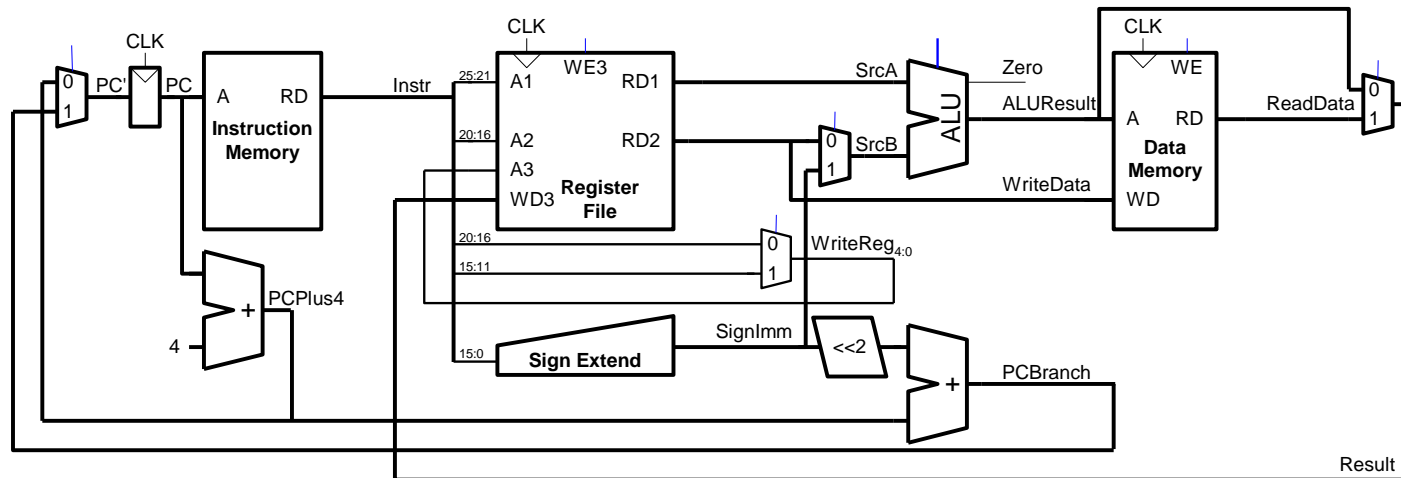
Pipelined



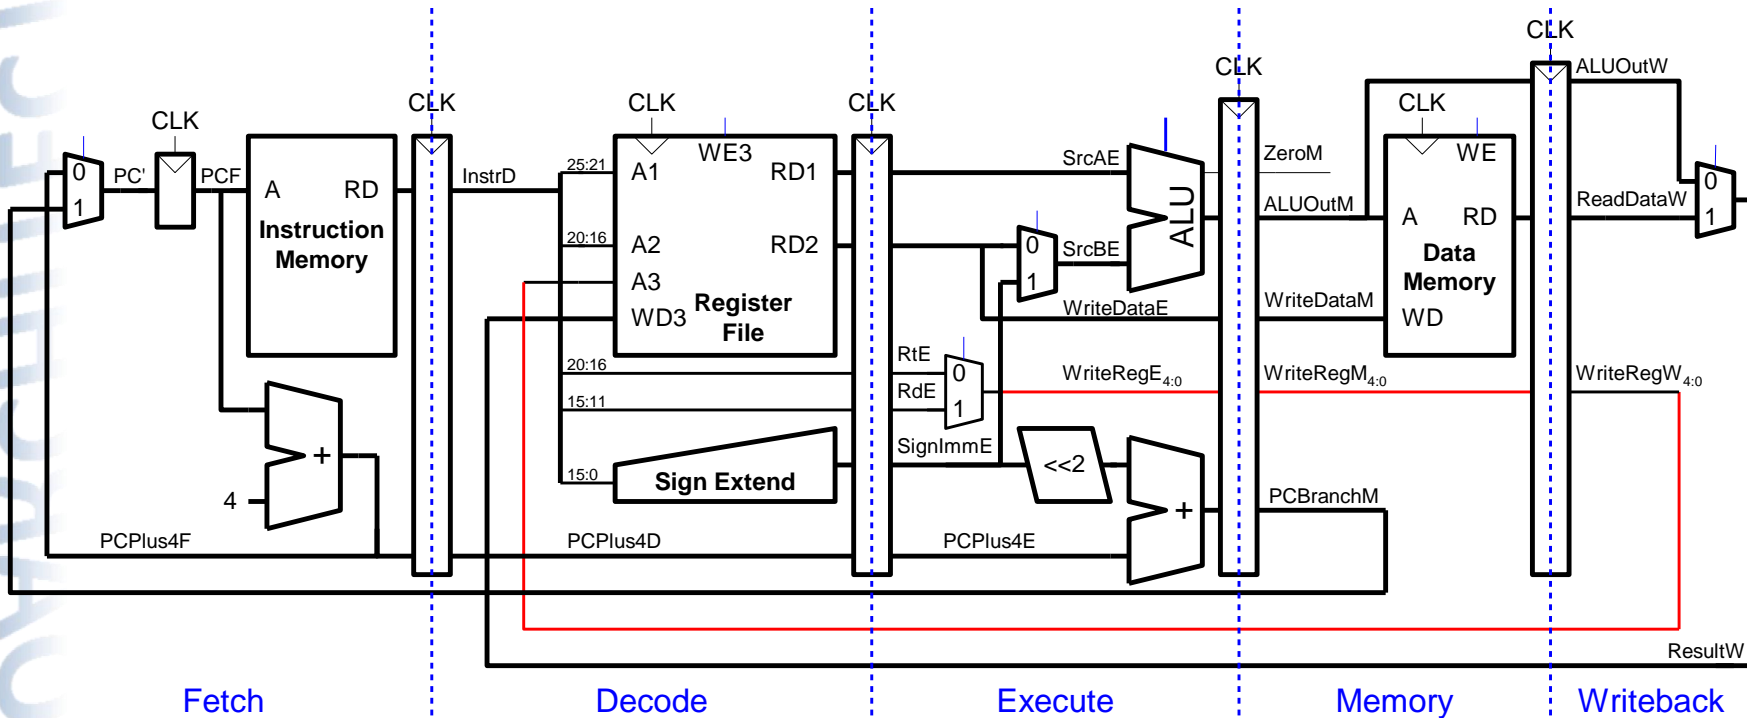
Pipelined Processor Abstraction



Single-Cycle & Pipelined Datapath

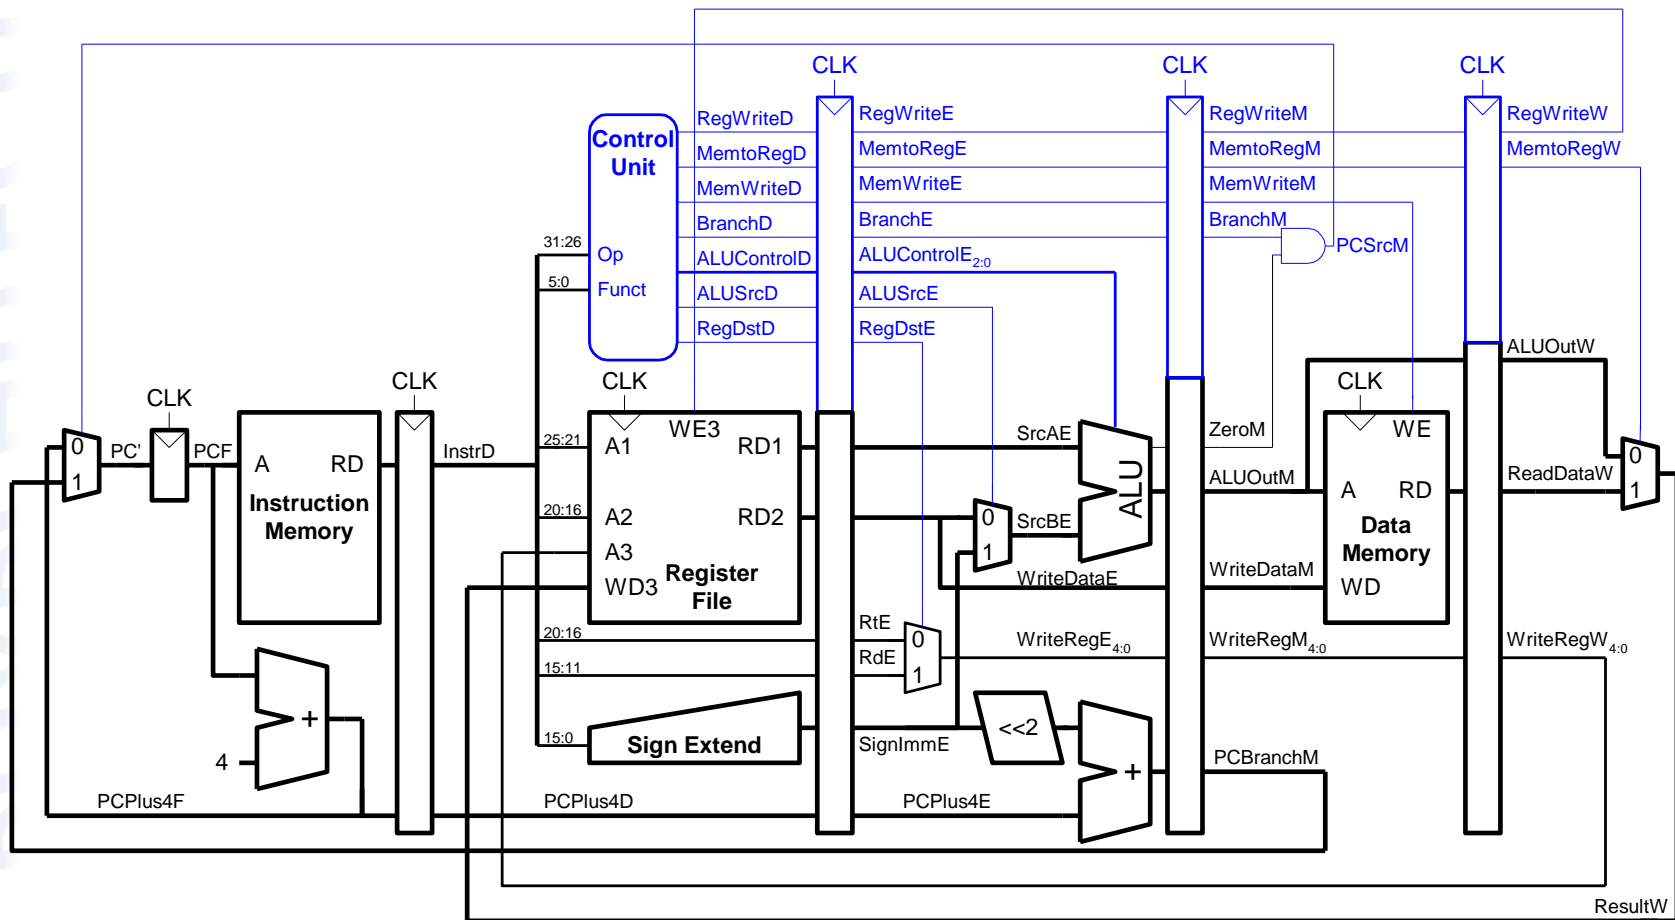


Corrected Pipelined Datapath



WriteReg must arrive at same time as *Result*

Pipelined Processor Control

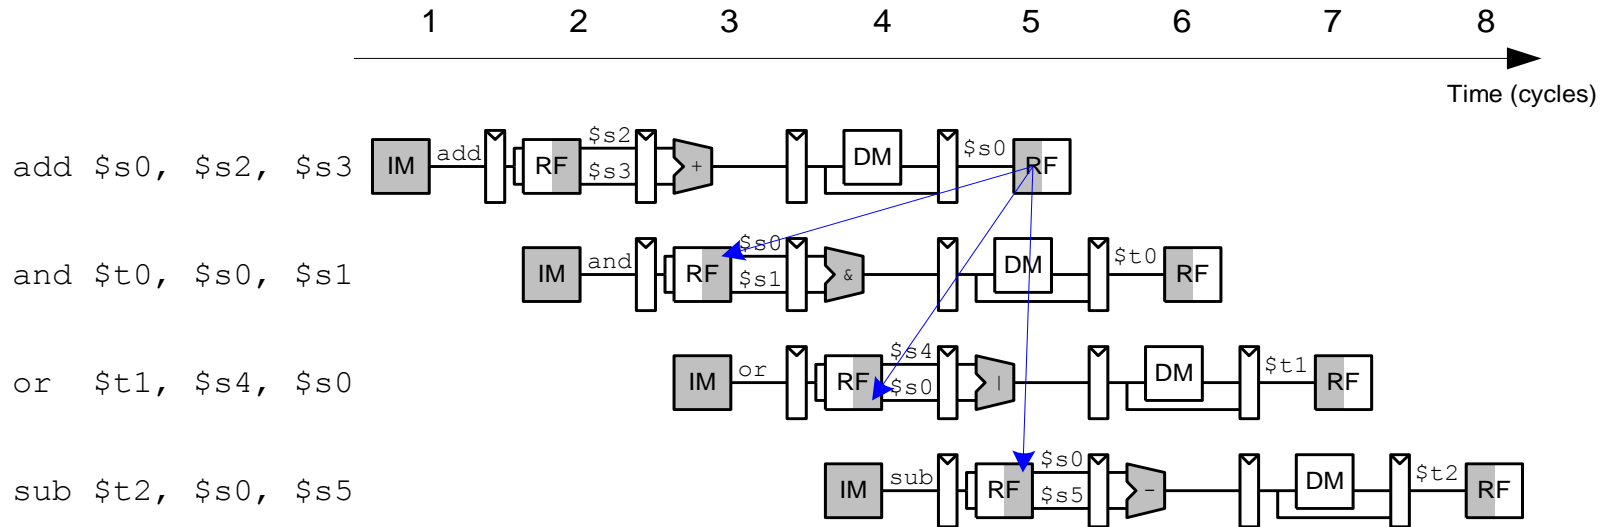


- Same control unit as single-cycle processor
- Control delayed to proper pipeline stage

Pipeline Hazards

- When an instruction depends on result from instruction that hasn't completed
- Types:
 - **Data hazard:** register value not yet written back to register file
 - **Control hazard:** next instruction not decided yet (caused by branches) or target address not calculated yet (jumps and branches)

Data Hazard



Handling Data Hazards

2 SW fixes

instructions that do nothing

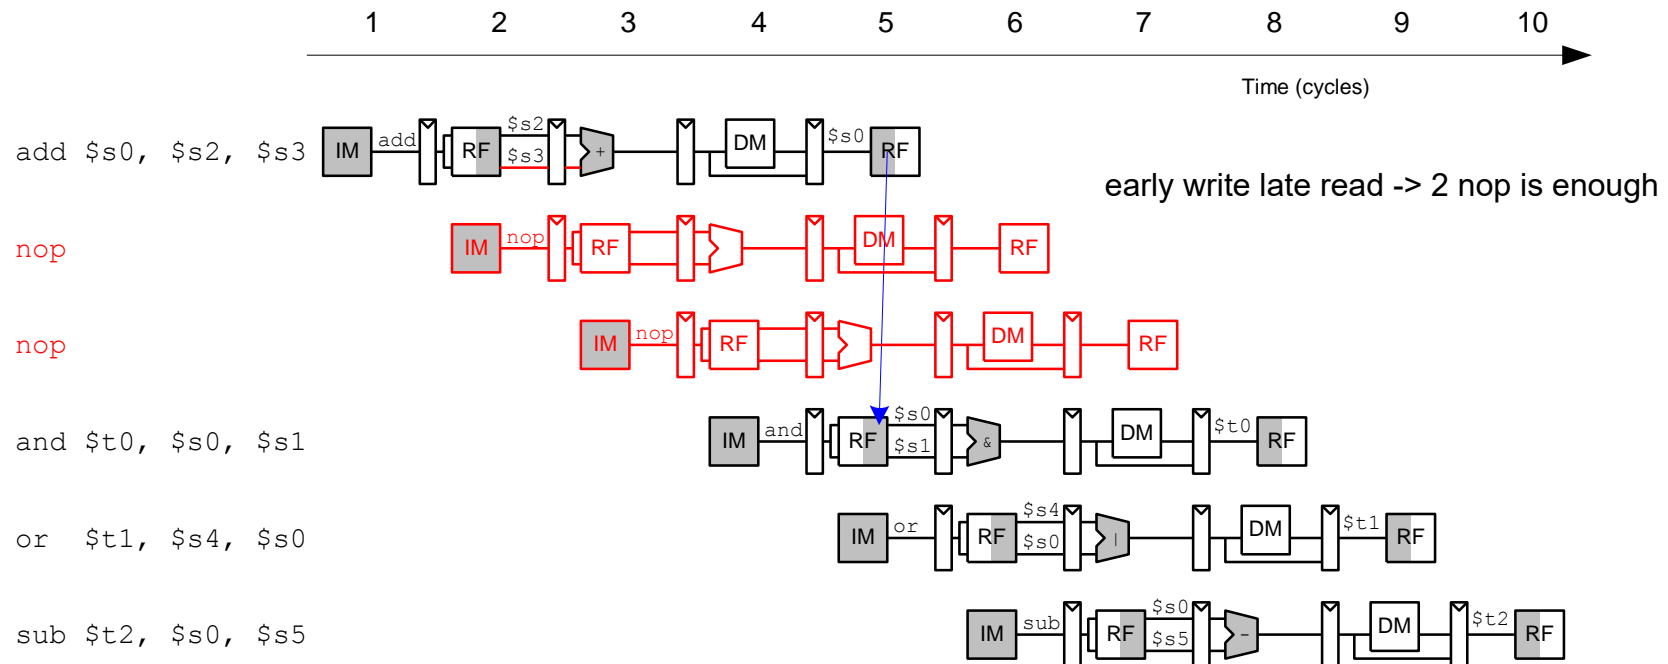
- Insert `nops` in code at compile time
- Rearrange code at compile time

2 HW fixes

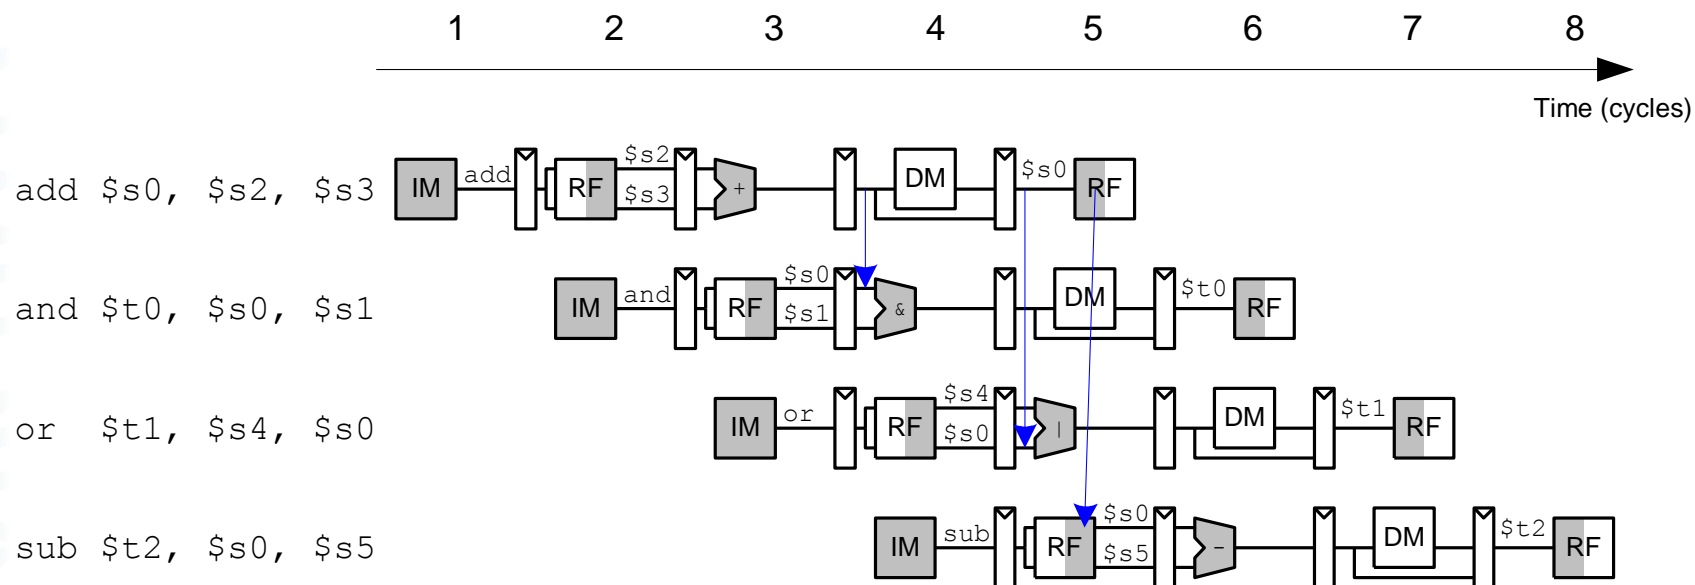
- Stall the processor at run time
- Forward data at run time

Compile-Time Hazard Elimination

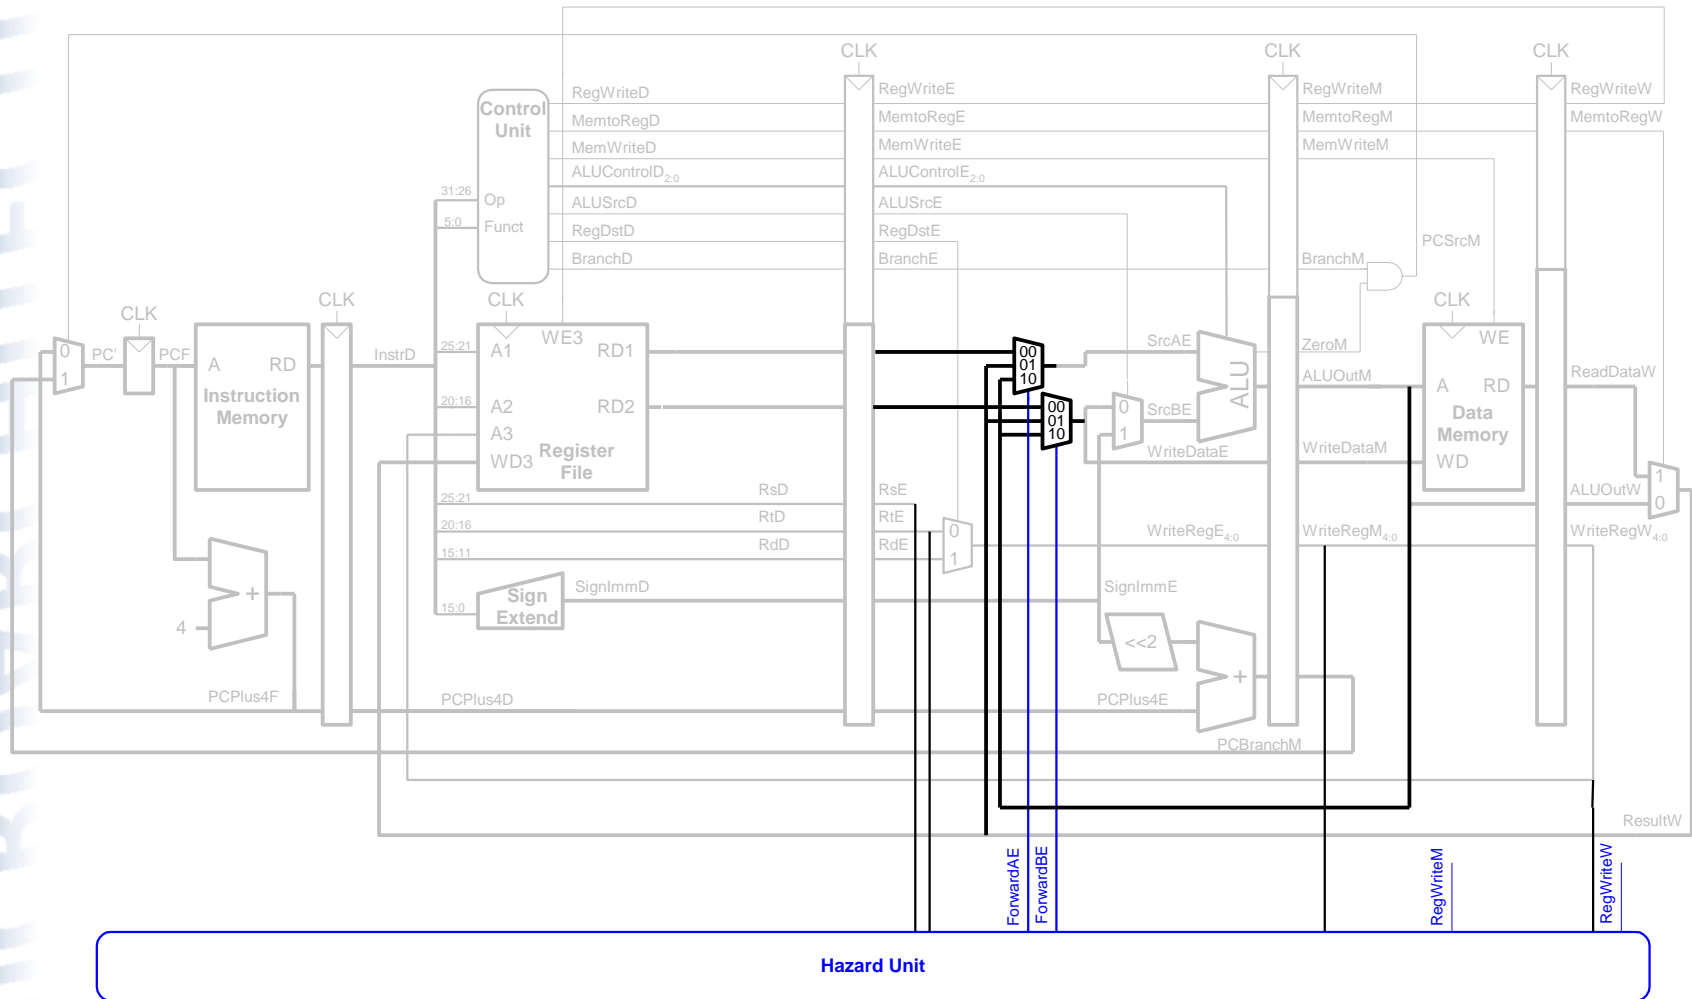
- Insert enough nops for result to be ready
- Or move independent useful instructions forward



Data Forwarding



Data Forwarding



Data Forwarding

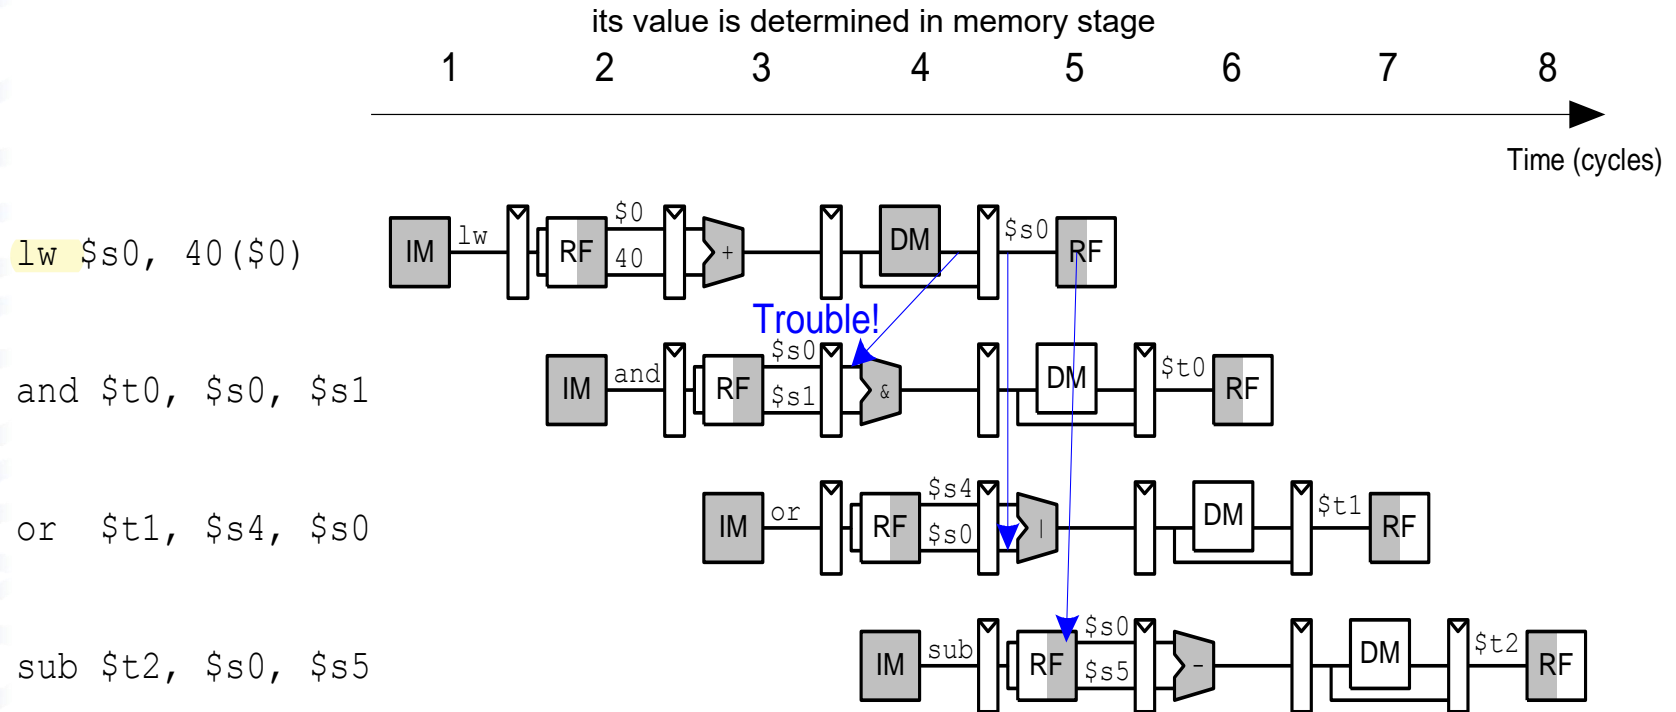
- Forward to Execute stage from either:
 - Memory stage or
 - Writeback stage
- Forwarding logic for *ForwardAE*:

```
if ((rsE != 0) AND (rsE == WriteRegM) AND RegWriteM)
    then ForwardAE = 10
else
    if ((rsE != 0) AND (rsE == WriteRegW) AND RegWriteW)
        then ForwardAE = 01
    else ForwardAE = 00
```

Forwarding logic for *ForwardBE* same, but replace *rsE* with *rtE*

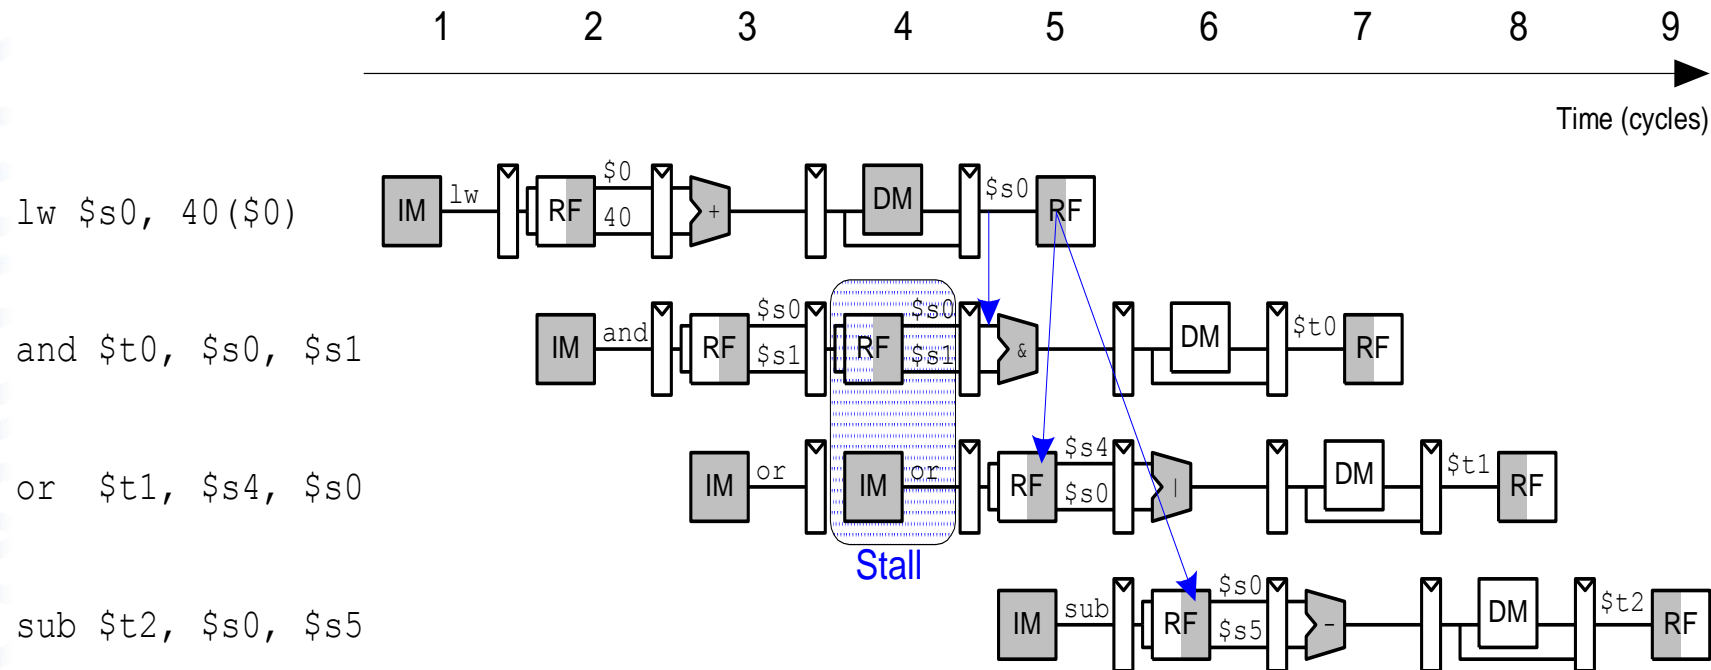
Stalling

Forwarding on a **load-use** hazard isn't possible!




Stalling

The HW solution is to stall the pipeline



Stalling Hardware



Stalling Logic

lwstall =

$((rsD == rtE) \text{ OR } (rtD == rtE)) \text{ AND } MemtoRegE$

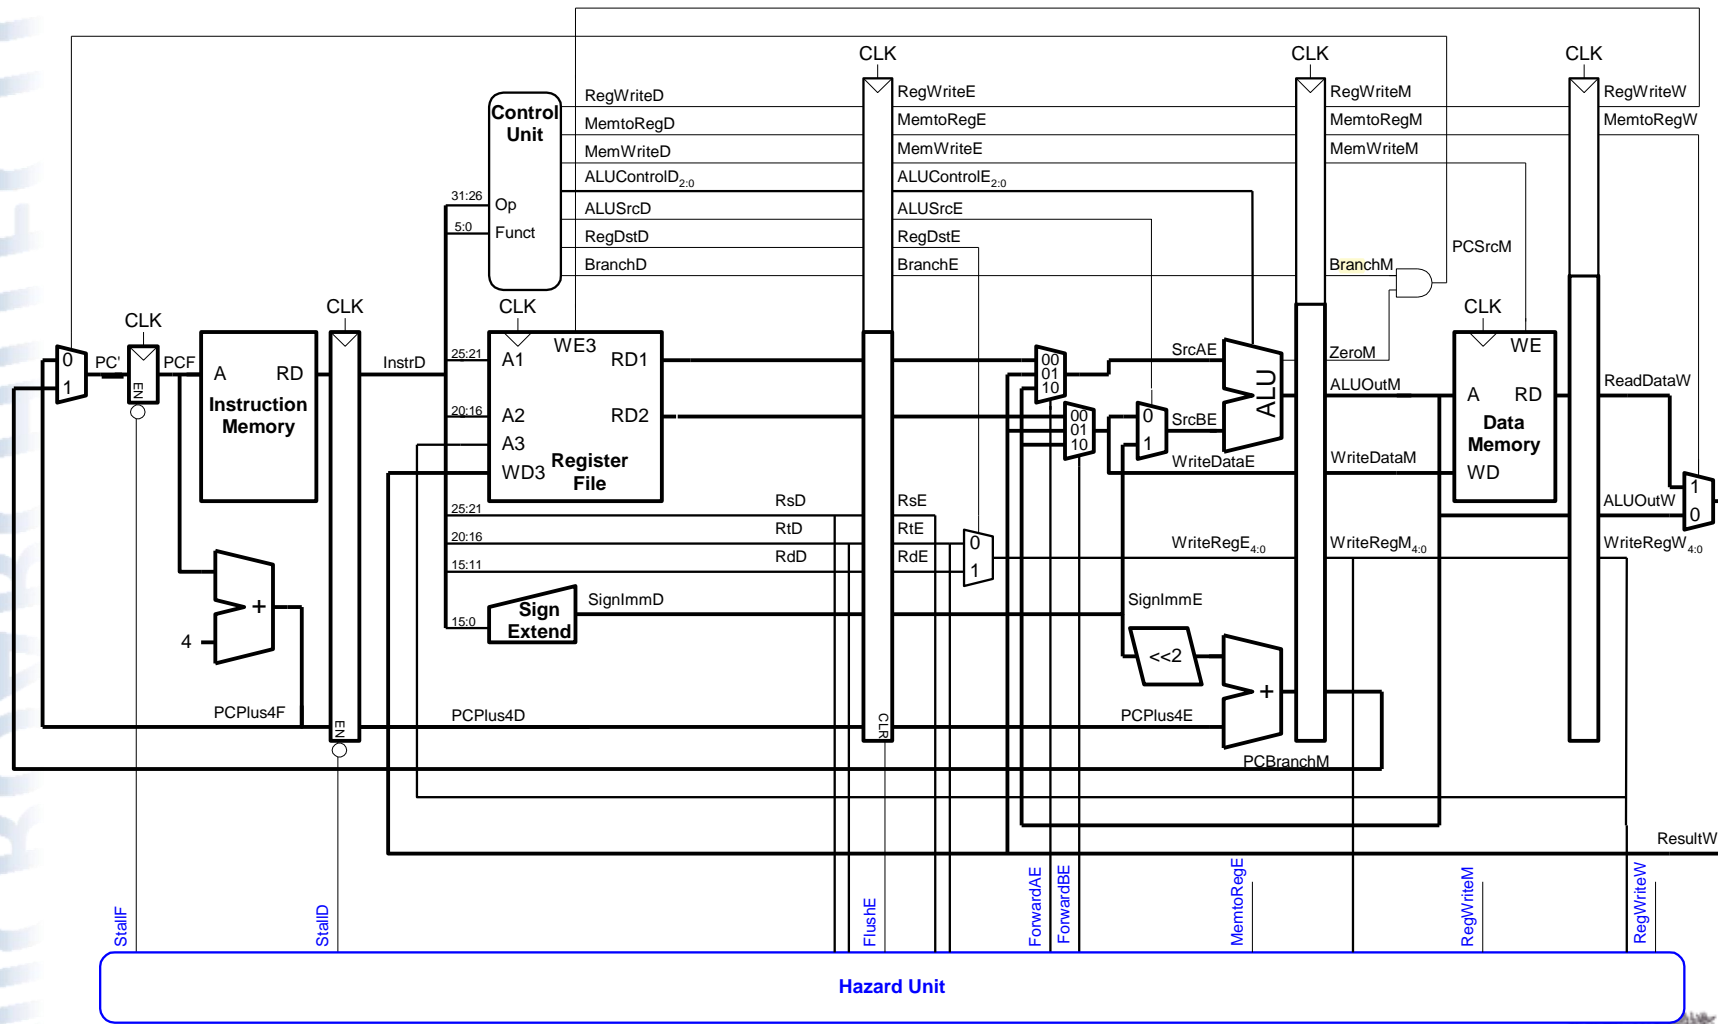
StallF = *StallD* = *FlushE* = *lwstall*

- By flushing the Execute stage, and stalling Fetch and Decode stages, the instruction flushed will simply be repeated in then next clock cycle, but this time with correct (forwarded) data!

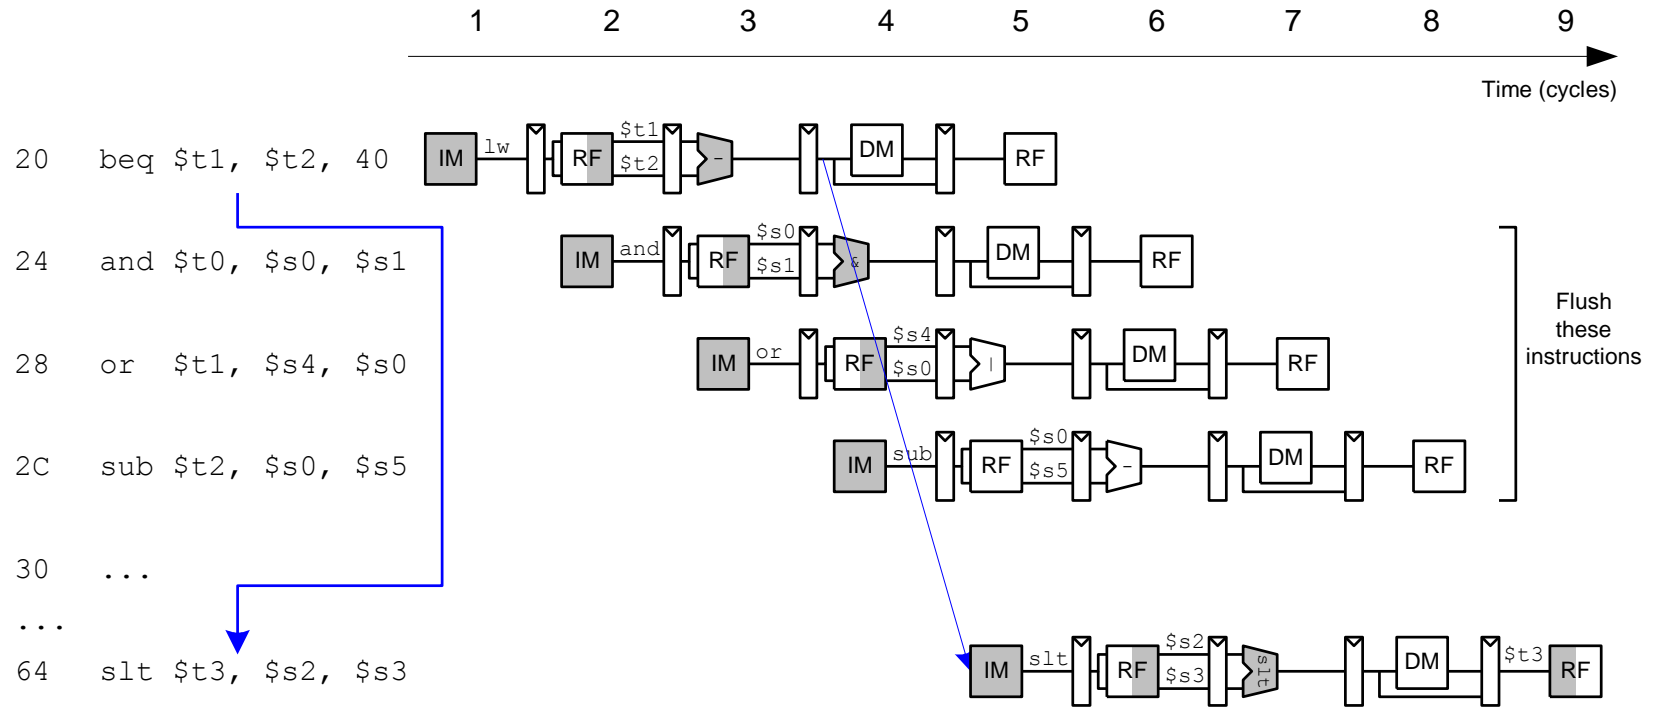
Control Hazards

- **beq:**
 - branch not determined until 4th stage of pipeline
 - Instructions after the branch are fetched *before* the branch occurs
 - These instructions must be flushed if branch happens
- **Branch misprediction penalty**
 - the # of instruction flushed, when branch is taken
 - may be reduced by determining branch earlier

Control Hazards: Original Pipeline

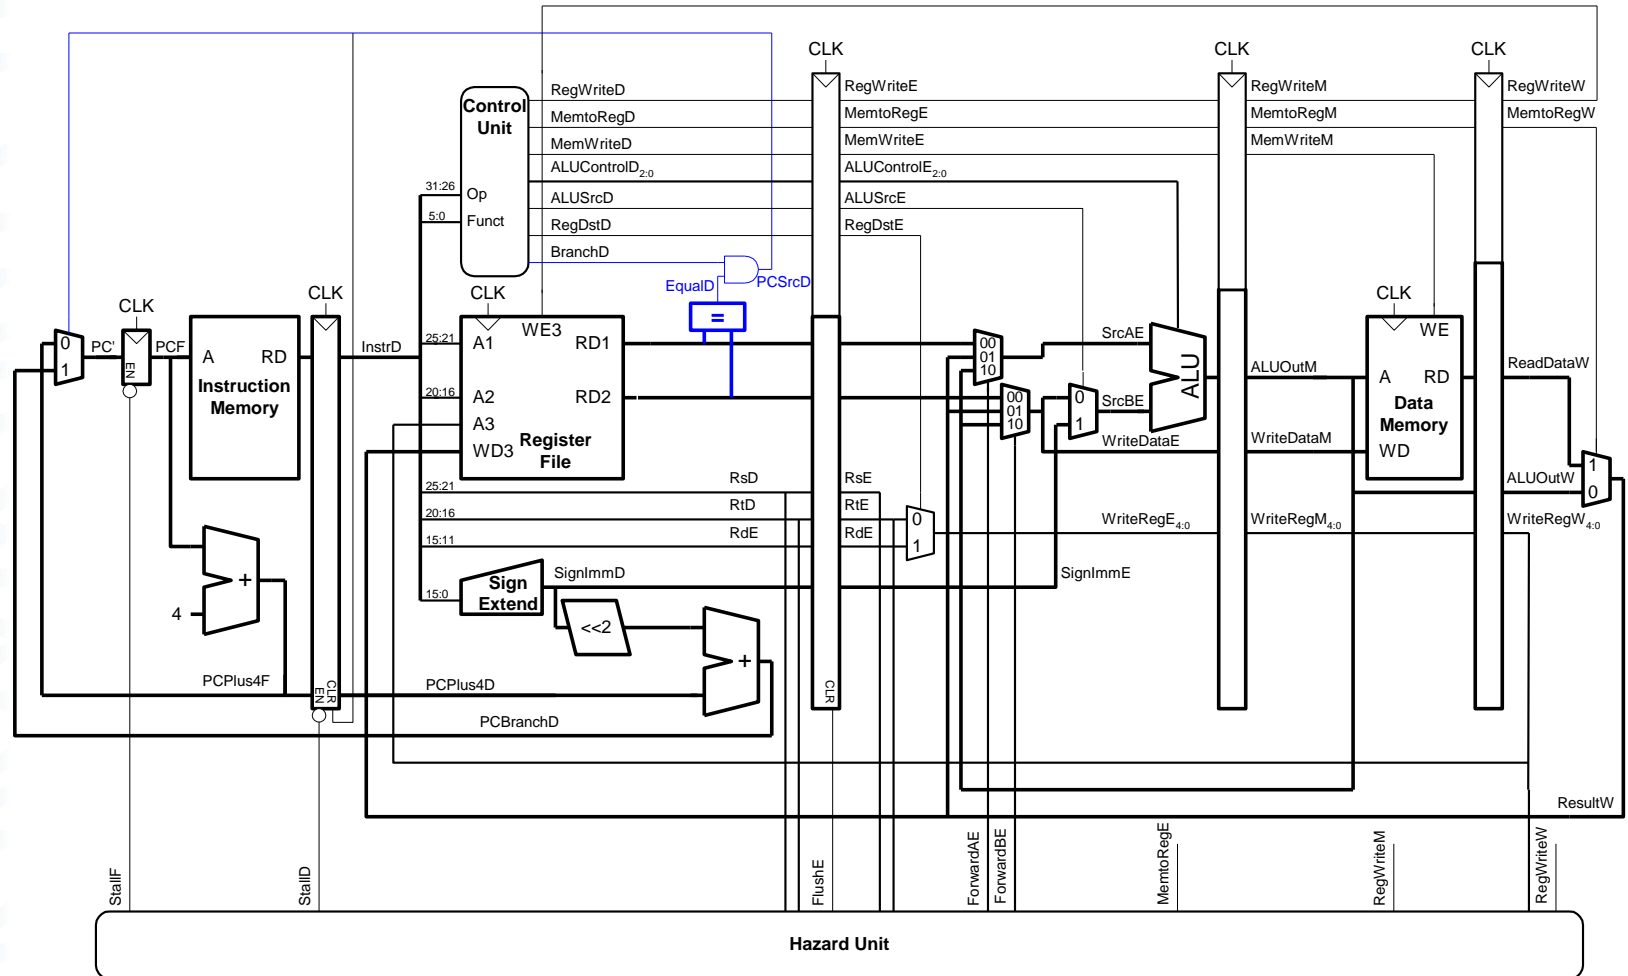


Control Hazards



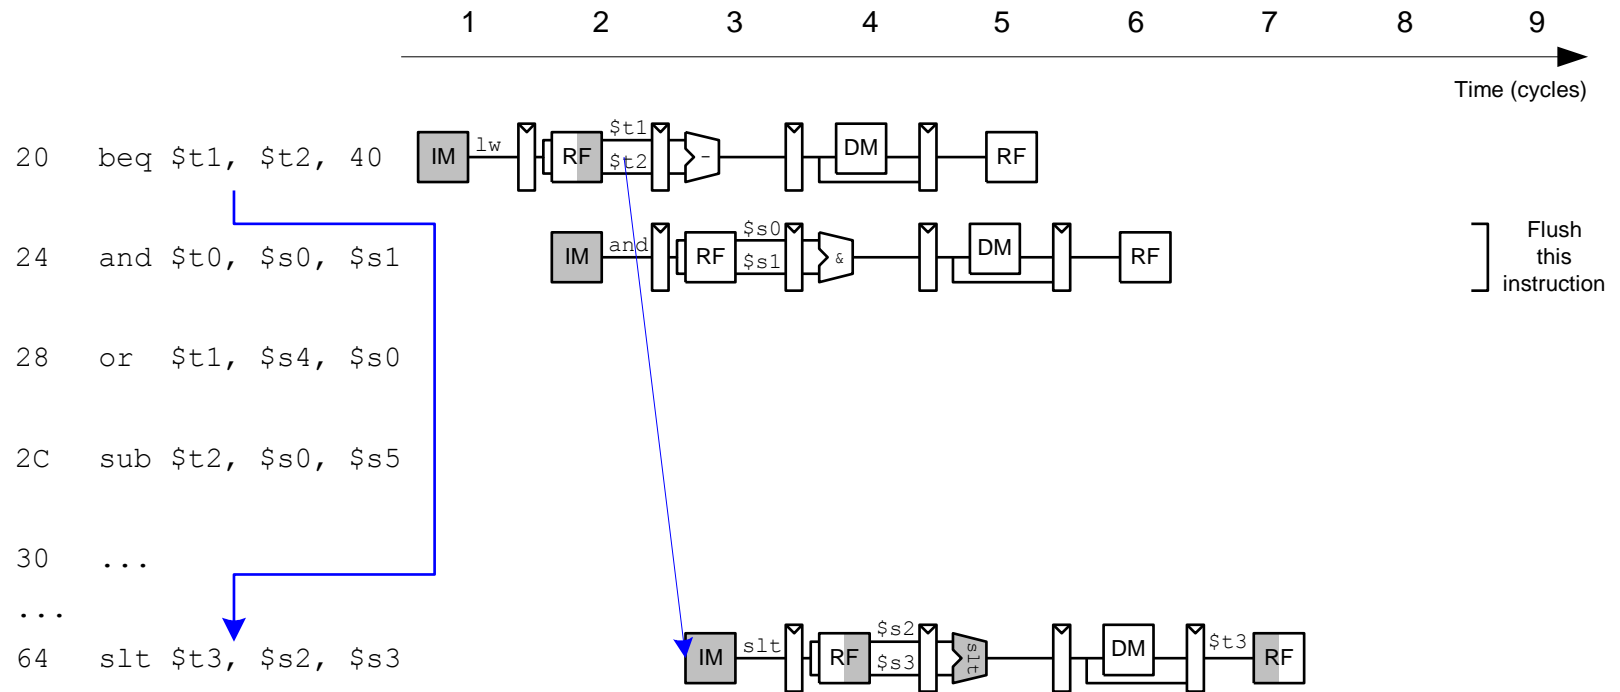
branch in memory stage --> 3 instruction branch penalty

Early Branch Resolution




But: introduced another data hazard in Decode stage!

Early Branch Resolution



branch in decode stage --> 1 instruction branch penalty

Forwarding to Early-branch HW



Control Forwarding & Stalling Logic

- **Forwarding logic:**

ForwardAD = (*rsD* != 0) AND (*rsD* == *WriteRegM*) AND *RegWriteM*

ForwardBD = (*rtD* != 0) AND (*rtD* == *WriteRegM*) AND *RegWriteM*

- **Stalling logic:**

branchstall = *BranchD* AND *RegWriteE* AND
(*WriteRegE* == *rsD* OR *WriteRegE* == *rtD*)

OR

BranchD AND *MemtoRegM* AND
(*WriteRegM* == *rsD* OR *WriteRegM* == *rtD*)

StallF = *StallD* = *FlushE* = (*lwstall* OR *branchstall*)



Branch Prediction

- Guess whether branch will be taken
 - Backward branches are usually taken (in bottom-tested loops)
 - Consider history to improve guess
- Good prediction significantly reduces fraction of branches requiring a flush
- Requires HW for history table, etc

Pipelined Performance Example

- SPECINT2000 benchmark:
 - 25% loads
 - 10% stores
 - 11% branches
 - 2% jumps
 - 52% R-type
- Suppose:
 - 40% of loads used by next instruction
 - 25% of branches mispredicted
 - All jumps flush next instruction (JTA not ready)
- **What is the average CPI?**

Calculation of Average CPI

- Average CPI is the **weighted average** of CPI_{lw} , CPI_{sw} , CPI_{beq} , CPI_j and CPI_{R-type}
- For pipeline processors, $CPI = 1 + \# \text{ of stall cycles}$

Load CPI = 1 when no stall, = 2 when load-use occurs (1 stall)

– $CPI_{lw} = 1(0.6) + 2(0.4) = 1.4$

– $CPI_{sw} = 1$

Branch CPI = 1 when no stall, = 2 when it mispredicts and stalls

– $CPI_{beq} = 1(0.75) + 2(0.25) = 1.25$

Jump CPI = 2 since it always requires 1 stall

– $CPI_j = 2$

– $CPI_{R-type} = 1$

$$\begin{aligned} \text{Average CPI} &= (0.25)(1.4) + (0.1)(1) + (0.11)(1.25) + (0.02)(2) + \\ &\quad (0.52)(1) = \mathbf{1.15} \end{aligned}$$



Pipelined Performance

- Pipelined processor critical path:

$$T_c = \max \{$$
$$t_{pcq} + t_{mem} + t_{setup}$$
$$2(t_{RFread} + t_{mux} + t_{eq} + t_{AND} + t_{mux} + t_{setup})$$
$$t_{pcq} + t_{mux} + t_{mux} + t_{ALU} + t_{setup}$$
$$t_{pcq} + t_{memwrite} + t_{setup}$$
$$2(t_{pcq} + t_{mux} + t_{RFwrite}) \}$$

Pipelined Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	t_{pcq_PC}	30
Register setup	t_{setup}	20
Multiplexer	t_{mux}	25
ALU	t_{ALU}	200
Memory read	t_{mem}	250
Register file read	t_{RFread}	150
Register file setup	$t_{RFsetup}$	20
Equality comparator	t_{eq}	40
AND gate	t_{AND}	15
Memory write	$T_{memwrite}$	220
Register file write	$t_{RFwrite}$	100 ps

$$\begin{aligned}
 T_c &= 2(t_{RFread} + t_{mux} + t_{eq} + t_{AND} + t_{mux} + t_{setup}) \\
 &= 2[150 + 25 + 40 + 15 + 25 + 20] \text{ ps} = \mathbf{550 \text{ ps}}
 \end{aligned}$$

Pipelined Performance Example

Program with IC = 100 billion instructions

$$\begin{aligned}\text{Execution Time} &= \text{IC} \times \text{CPI} \times T_c \\ &= (100 \times 10^9)(1.15)(550 \times 10^{-12}) \\ &= \mathbf{63 \text{ seconds}}\end{aligned}$$

Processor Performance Comparison

Processor	Execution Time (seconds)	Speedup (single-cycle as baseline)
Single-cycle	92.5	1
Multicycle	133	0.70
Pipelined	63	1.47

Advanced Microarchitecture

- Deep Pipelining
- Branch Prediction
- Superscalar Processors
- Out of Order Processors
- Register Renaming
- SIMD
- Multithreading
- Multiprocessors

Deep Pipelining

- 10-20 stages typical
- Number of stages limited by:
 - Pipeline hazards
 - Sequencing overhead
 - Power
 - Cost

Branch Prediction

- Ideal pipelined processor: $CPI = 1$
- Branch misprediction increases CPI
- **Static branch prediction:**
 - Check direction of branch (forward or backward)
 - If backward, predict taken
 - Else, predict not taken
- **Dynamic branch prediction:**
 - Keep history of last (several hundred) branches in *branch target buffer*, record:
 - Branch destination
 - Whether branch was taken

Branch Prediction Example


```
add    $s1, $0, $0          # sum = 0
add    $s0, $0, $0          # i   = 0
addi   $t0, $0, 10          # $t0 = 10
for:
    beq  $s0, $t0, done      # if i == 10, branch
    add  $s1, $s1, $s0       # sum = sum + i
    addi $s0, $s0, 1         # increment i
    j    for
done:
```

1-Bit Branch Predictor

- Remembers whether branch was taken the last time and does the same thing
- Mispredicts first and last branch of loop

TAKEN -> NOT TAKEN X9 -> TAKEN

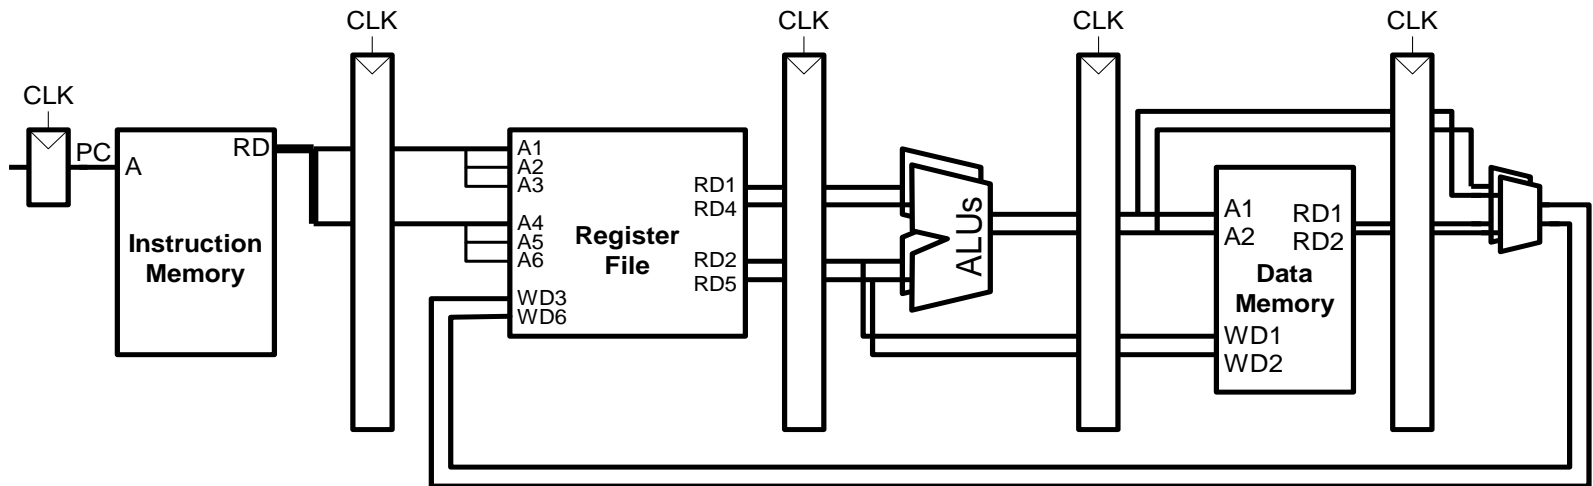
2-Bit Branch Predictor



Only mispredicts the last branch of the loop

Superscalar

- Multiple copies of datapath execute multiple instructions at once
- Dependencies make it tricky to issue multiple instructions at once

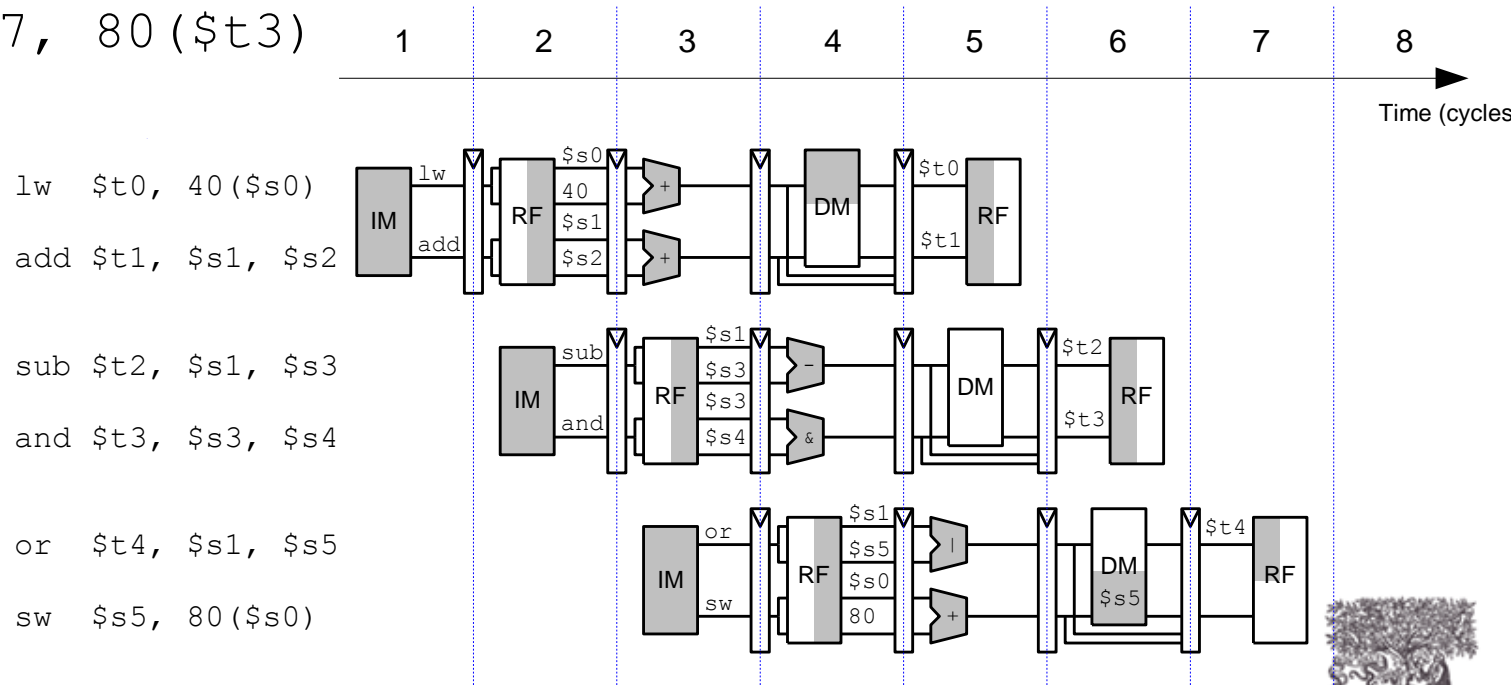


Superscalar Example

```
lw    $t0, 40($s0)
add   $t1, $t0, $s1
sub   $t0, $s2, $s3
and   $t2, $s4, $t0
or    $t3, $s5, $s6
sw    $s7, 80($t3)
```

Ideal IPC: 2

Actual IPC: 2



Superscalar with Dependencies

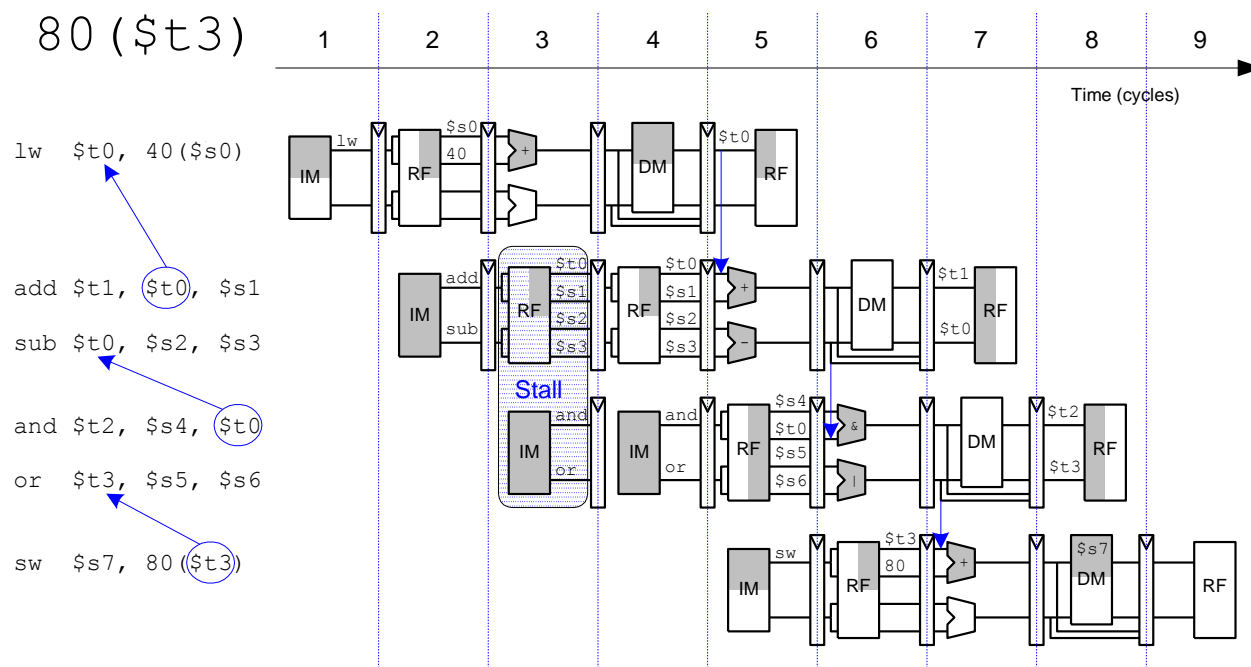
```
lw    $t0, 40($s0)
add   $t1, $t0, $s1
sub   $t0, $s2, $s3
and   $t2, $s4, $t0
or    $t3, $s5, $s6
sw    $s7, 80($t3)
```

Ideal IPC:

2

Actual IPC:

6/5 = 1.17



Out of Order Processor

- Looks ahead across multiple instructions
- Issues as many instructions as possible at once
- Issues instructions out of order (as long as no dependencies)
- **Dependencies:**
 - **RAW** (read after write): one instruction writes, later instruction reads a register
 - **WAR** (write after read): one instruction reads, later instruction writes a register
 - **WAW** (write after write): one instruction writes, later instruction writes a register

Out of Order Processor

- **Instruction level parallelism (ILP):** number of instruction that can be issued simultaneously (average < 3)
- **Scoreboard:** table that keeps track of:
 - Instructions waiting to issue
 - Available functional units
 - Dependencies

Out of Order Processor Example

```
lw    $t0, 40($s0)
add   $t1, $t0, $s1
sub   $t0, $s2, $s3
and   $t2, $s4, $t0
or    $t3, $s5, $s6
sw    $s7, 80($t3)
```

Ideal IPC:

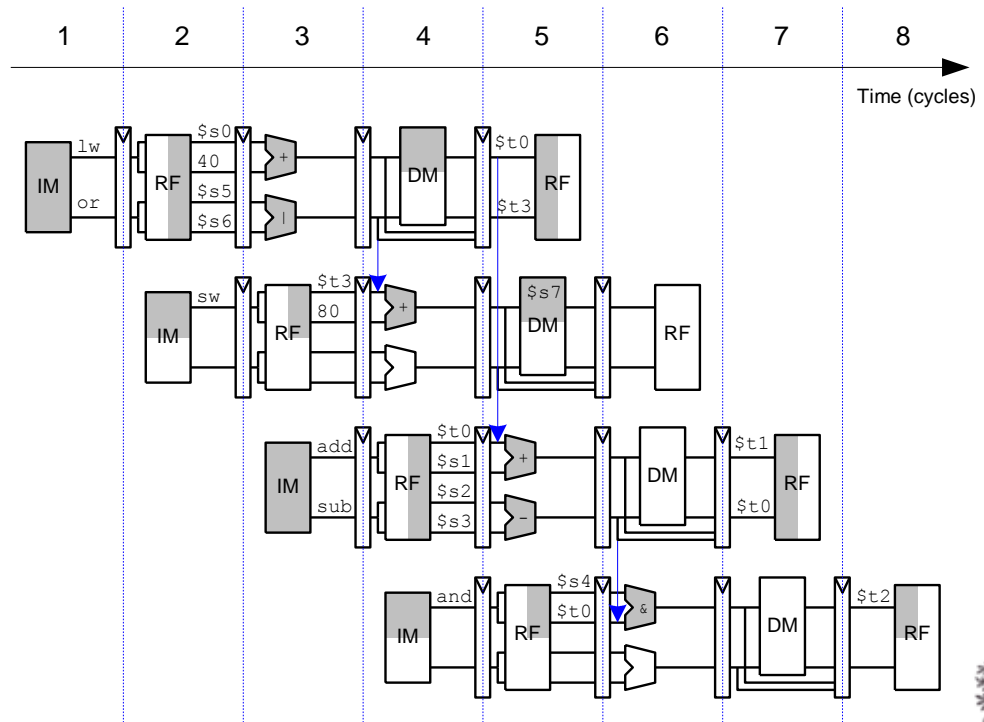
2

Actual IPC:

$6/4 = 1.5$

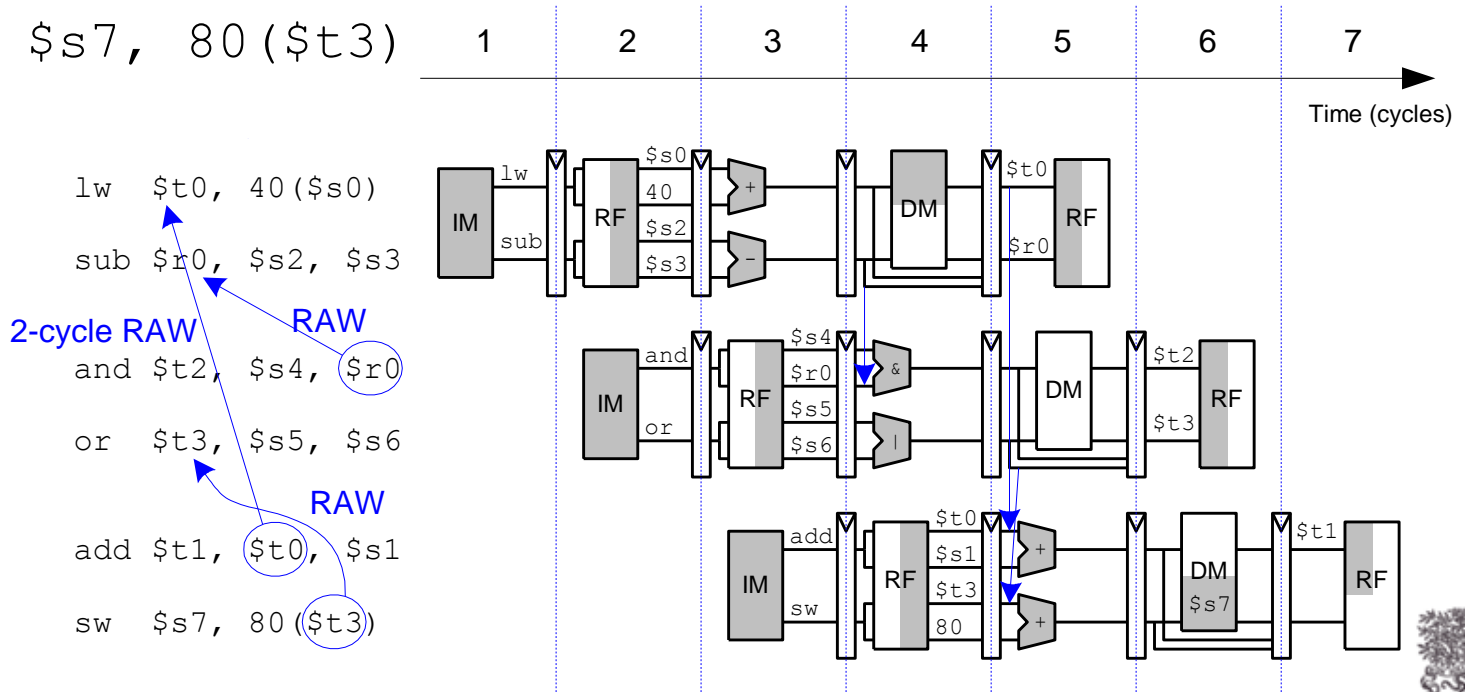
lw \$t0, 40(\$s0)
 or \$t3, \$s5, \$s6
 sw \$s7, 80(\$t3)
 add \$t1, \$t0, \$s1
 sub \$t0, \$s2, \$s3
 and \$t2, \$s4, \$t0

RAW
 two cycle latency
 between load and
 use of \$t0
 RAW
 WAR
 RAW



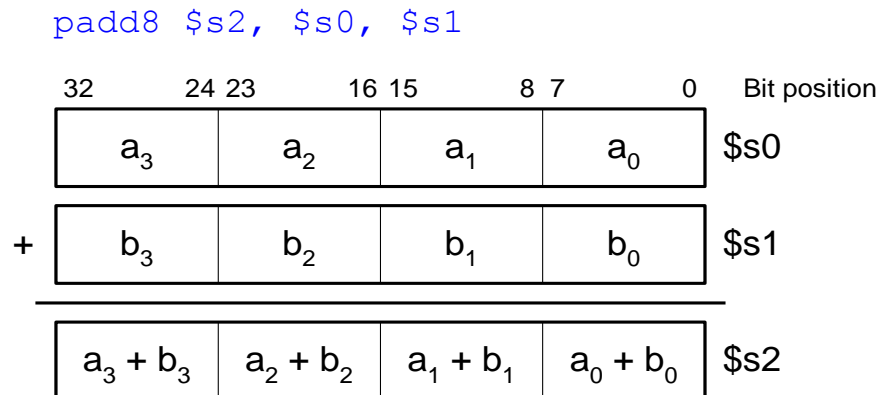
```
lw    $t0, 40($s0)
add   $t1, $t0, $s1
sub   $t0, $s2, $s3
and   $t2, $s4, $t0
or    $t3, $s5, $s6
sw    $s7, 80($t3)
```

2

$$6/3 = 2$$


SIMD

- Single Instruction Multiple Data (SIMD)
 - Single instruction acts on multiple pieces of data at once
 - Common application: graphics
 - Perform short arithmetic operations (also called *packed arithmetic*)
- For example, add four 8-bit elements



Advanced Architecture Techniques

- **Multithreading**
 - Wordprocessor: thread for typing, spell checking, printing
- **Multiprocessors**
 - Multiple processors (cores) on a single chip

Threading: Definitions

- **Process:** program running on a computer
 - Multiple processes can run at once: e.g., surfing Web, playing music, writing a paper
- **Thread:** part of a program
 - Each process has multiple threads: e.g., a word processor may have threads for typing, spell checking, printing

Threads in Conventional Processor

- One thread runs at once
- When one thread stalls (for example, waiting for memory):
 - Architectural state of that thread stored
 - Architectural state of waiting thread loaded into processor and it runs
 - Called **context switching**
- Appears to user like all threads running simultaneously

Multithreading

- Multiple copies of architectural state
- Multiple threads **active** at once:
 - When one thread stalls, another runs immediately
 - If one thread can't keep all execution units busy, another thread can use them
- Does not increase instruction-level parallelism (ILP) of single thread, but increases throughput

Intel calls this “hyperthreading”

Multiprocessors

- Multiple processors (cores) with a method of communication between them
- Types:
 - **Homogeneous:** multiple cores with shared memory
 - **Heterogeneous:** separate cores for different tasks (for example, DSP and CPU in cell phone)
 - **Clusters:** each core has own memory system